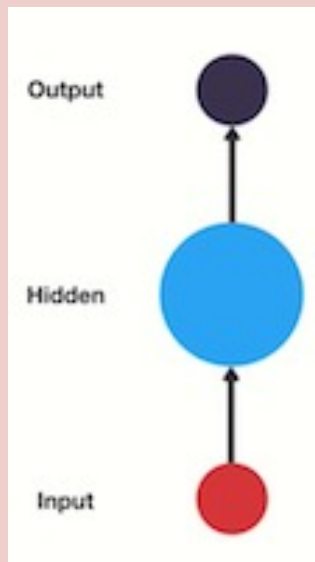


# **Seq2seq model (Encoder-Decoder model)**


Michel RIVEILL  
michel.riveill@univ-cotedazur.fr

# Remember

From

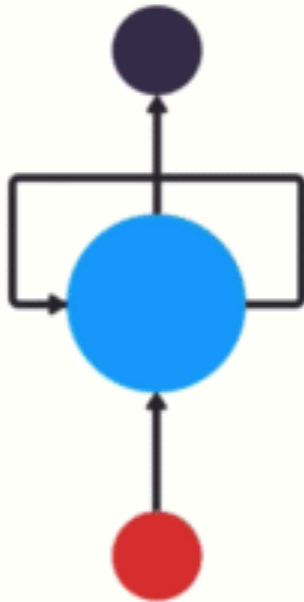


To

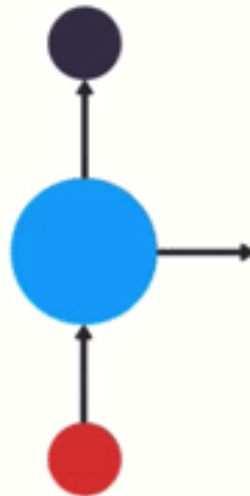


# Remember

From



To




# RNN in action



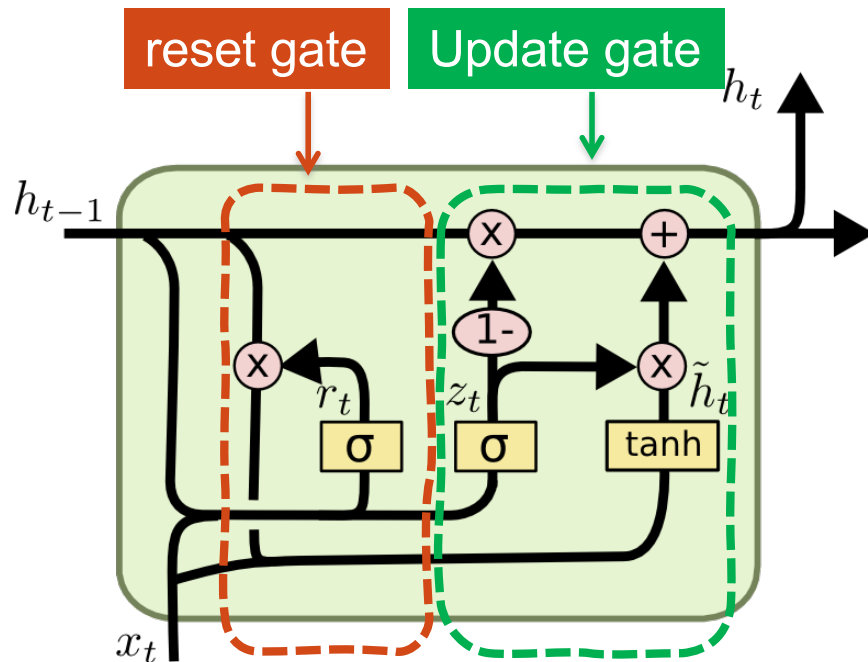
# LSTM cell

- ▶ Cell made up of three "gates": these are calculation zones which regulate the flow of information (by carrying out specific actions).
  - ▶ Forget gate (porte d'oubli)
  - ▶ Input gate (porte d'entrée)
  - ▶ Output gate (porte de sortie)
- ▶ Cell state (état de la cellule)
  - ▶ Like residual
- ▶ Hidden state (état caché)



# GRU – gated recurrent unit

- ▶ GRU = a light LSTM Cell



$$z_t = \sigma (W_z \cdot [h_{t-1}, x_t])$$

$$r_t = \sigma (W_r \cdot [h_{t-1}, x_t])$$


$$\tilde{h}_t = \tanh (W \cdot [r_t * h_{t-1}, x_t])$$

$$h_t = (1 - z_t) * h_{t-1} + z_t * \tilde{h}_t$$

- It combines the **forget** and **input** into a single **update gate**.
  - It also **merges** the **cell state** and **hidden state**.
- This is simpler than LSTM.

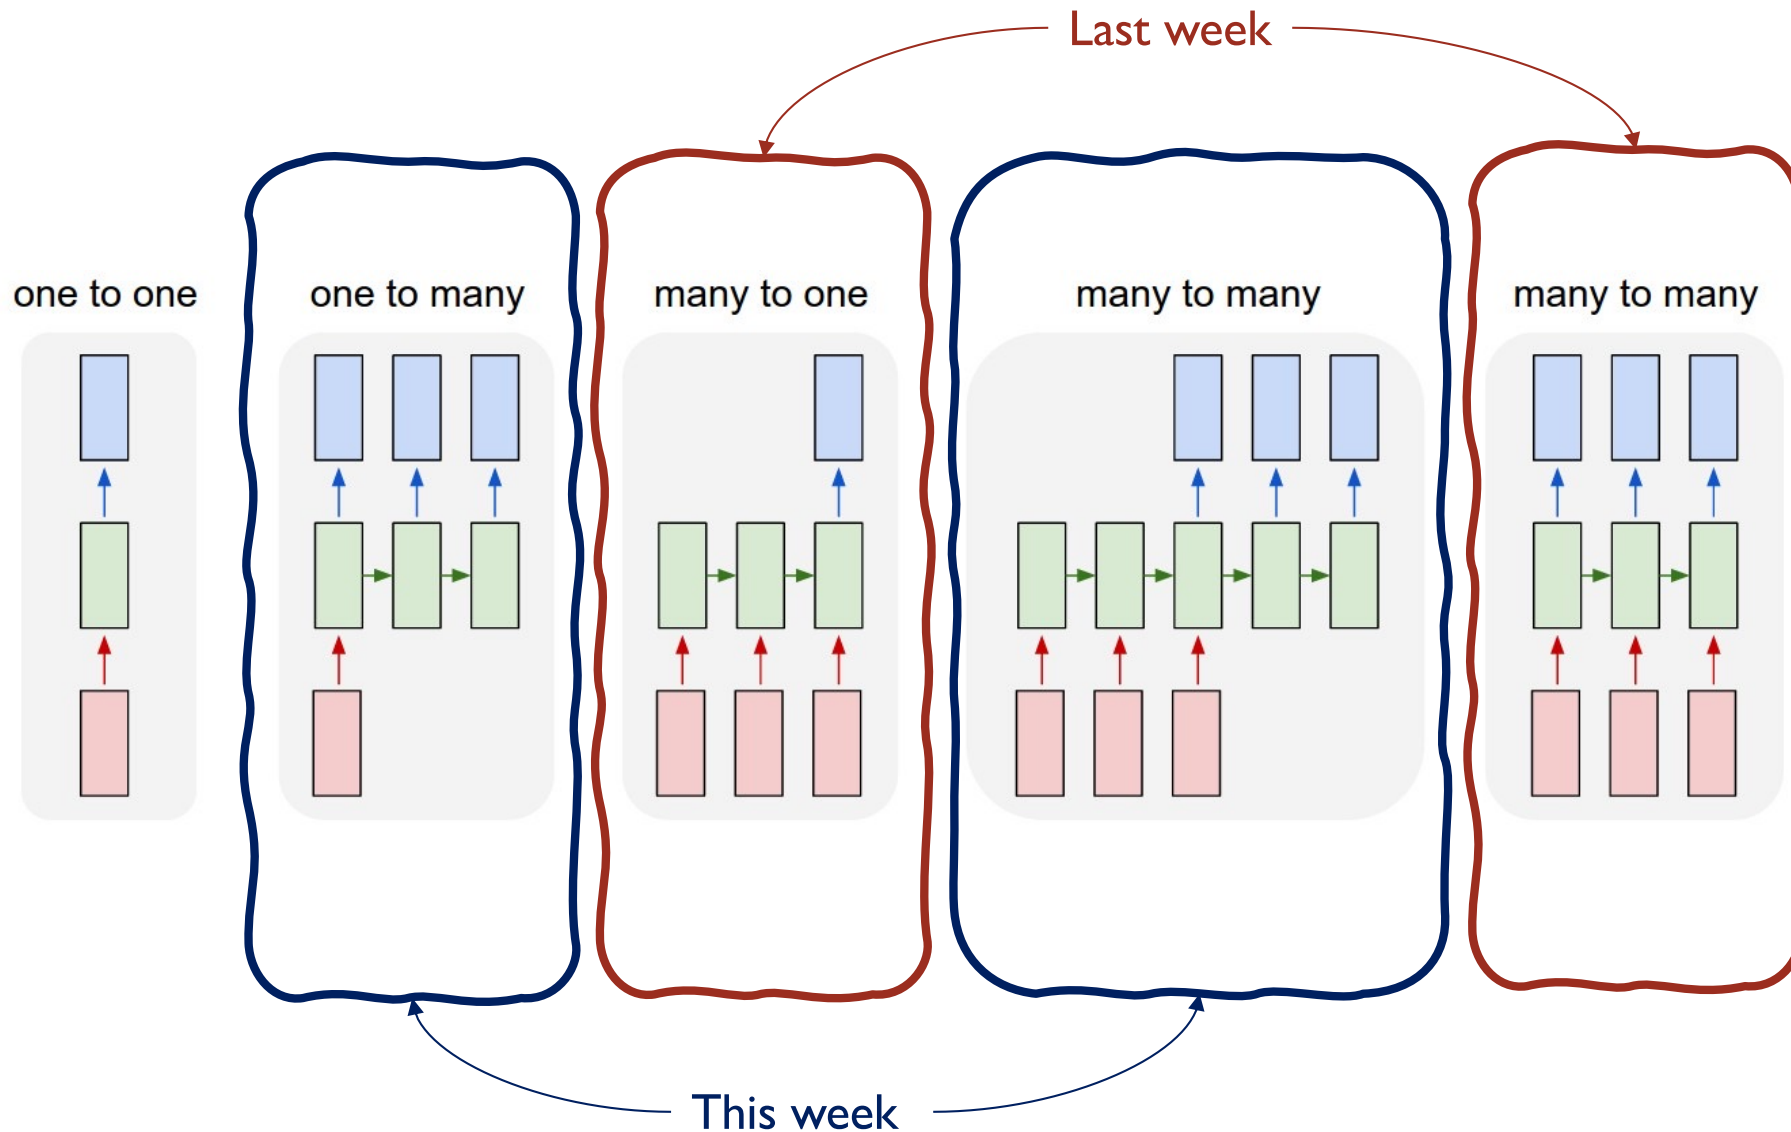
# Bi-directional RNNs

- ▶ RNNs can process the input sequence in forward and in the reverse direction



Popular in speech recognition used also with text

# Different RNN architecture

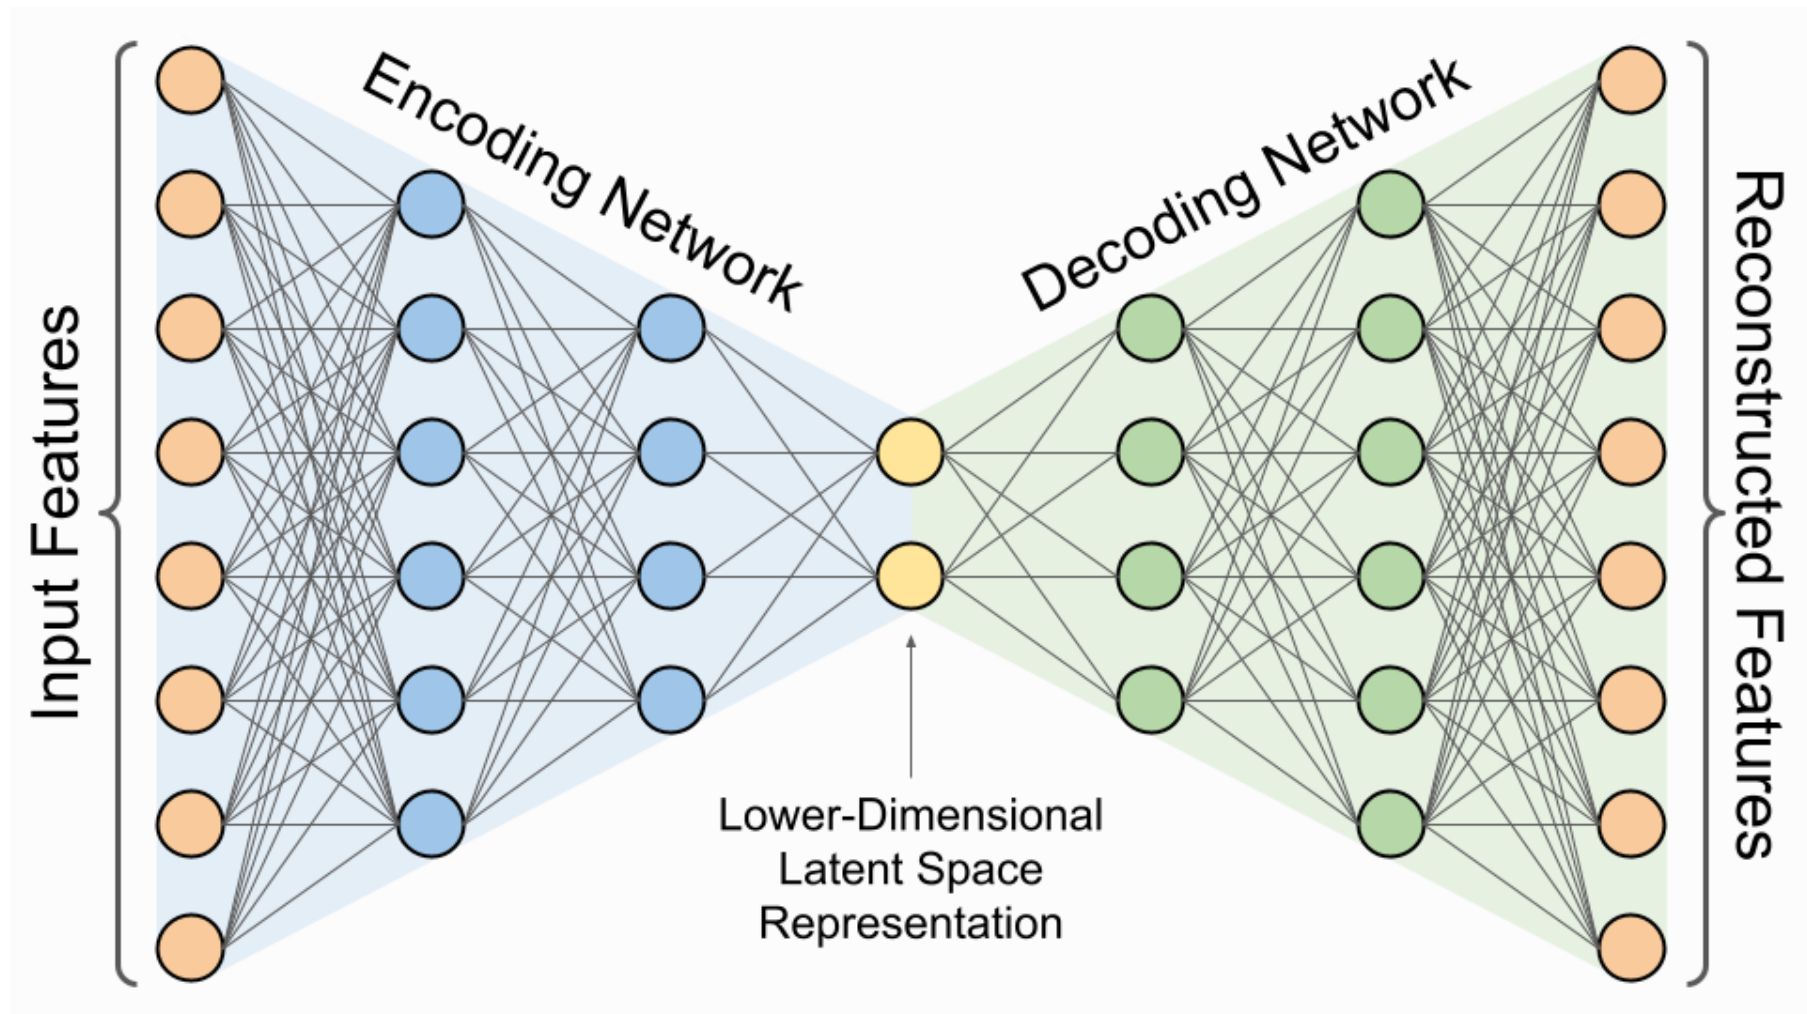




# **Seq2seq model (Encoder-Decoder model)**

$\text{len(input)} \neq \text{len(output)}$   
And generally, the two lengths are unknown

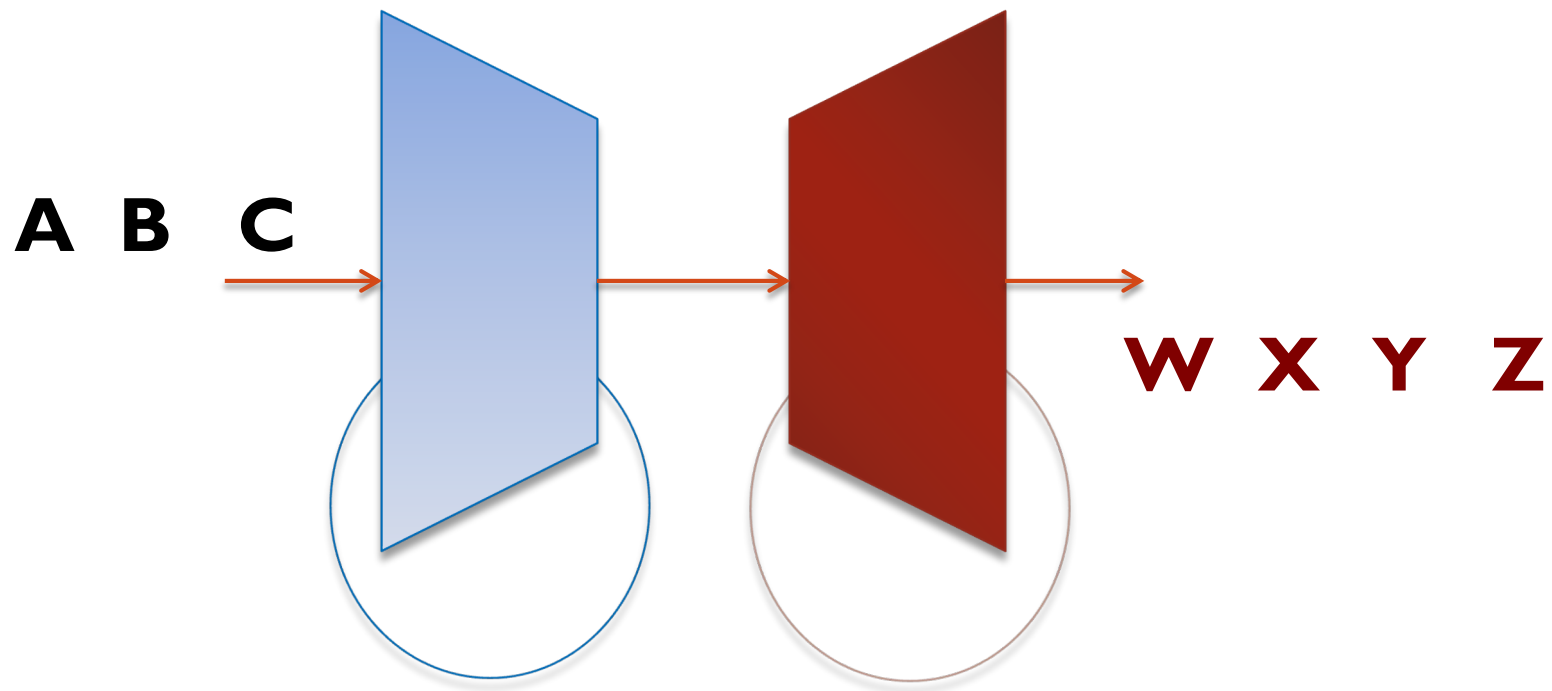
# Auto-encoder architecture



# Seq2seq – original approach

---

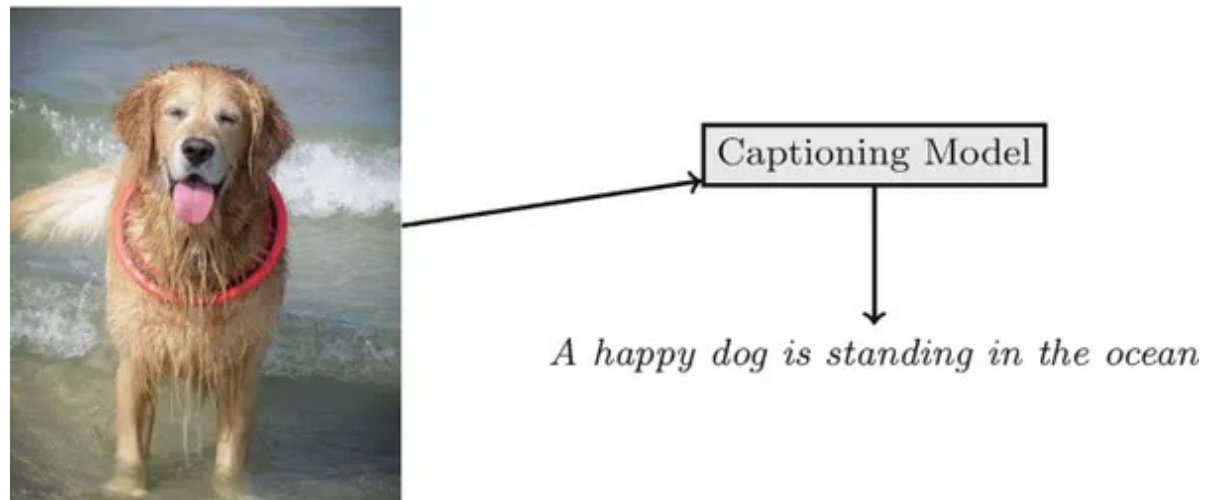
- ▶ Extend encoder-decoder architecture
  - ▶ to a **sequence data**
  - ▶ in order to develop an architecture capable of generating contextually appropriate, arbitrary length, output sequences



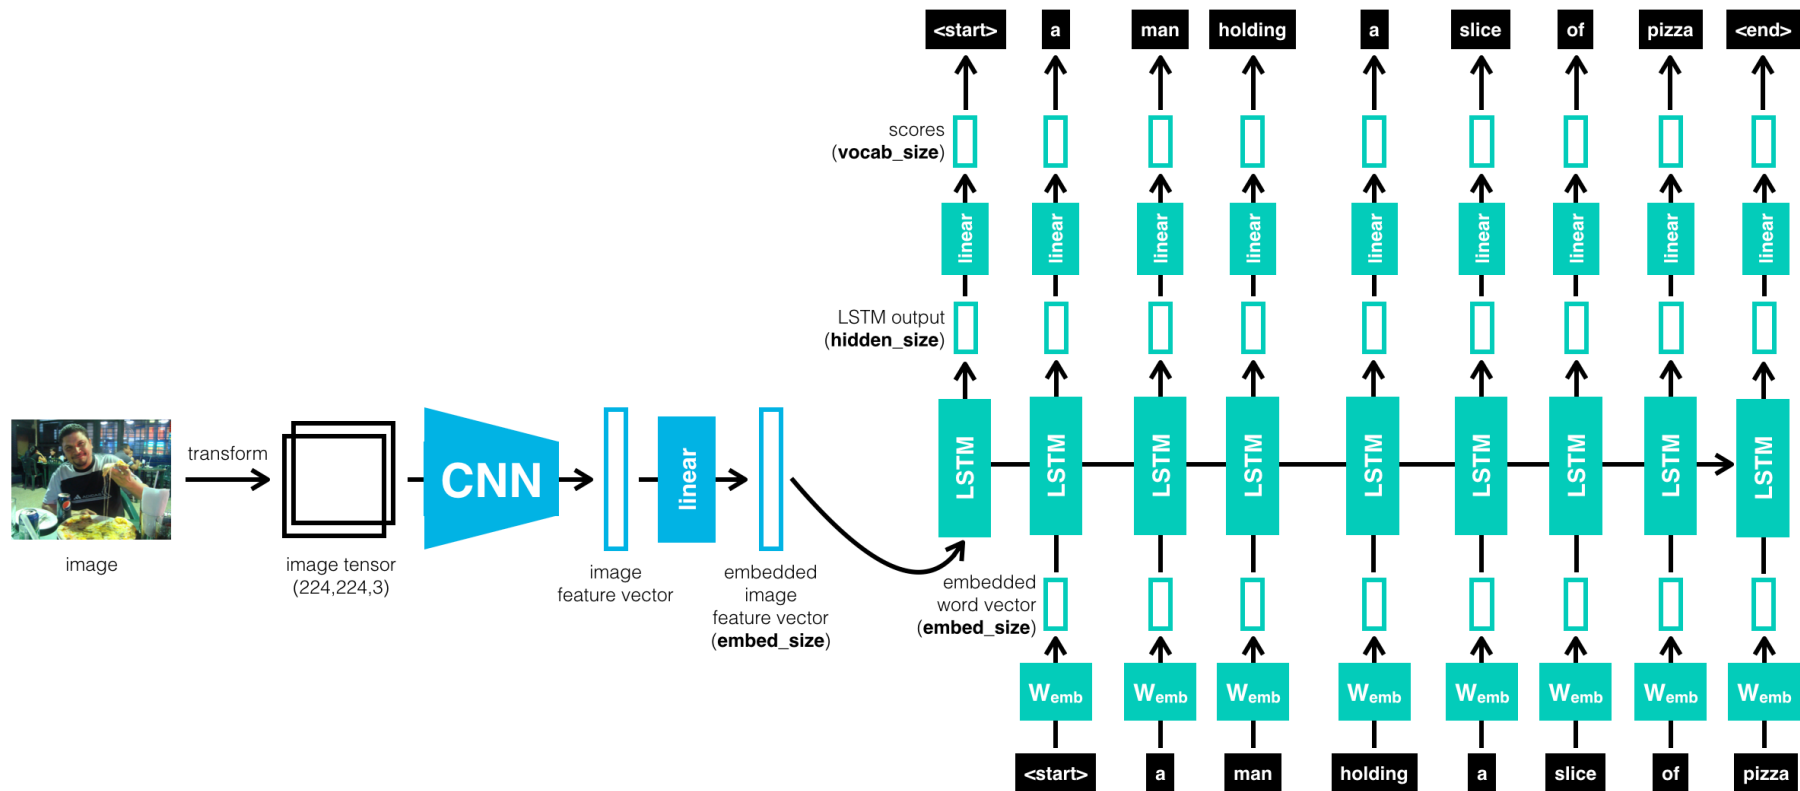
# Seq2Seq model

---

- ▶ Applications with text
  - ▶ Machine translation
  - ▶ Text summarization
  - ▶ Question answering
  - ▶ Dialogue modeling
- ▶ But also
  - ▶ Forecasting
  - ▶ Image captioning

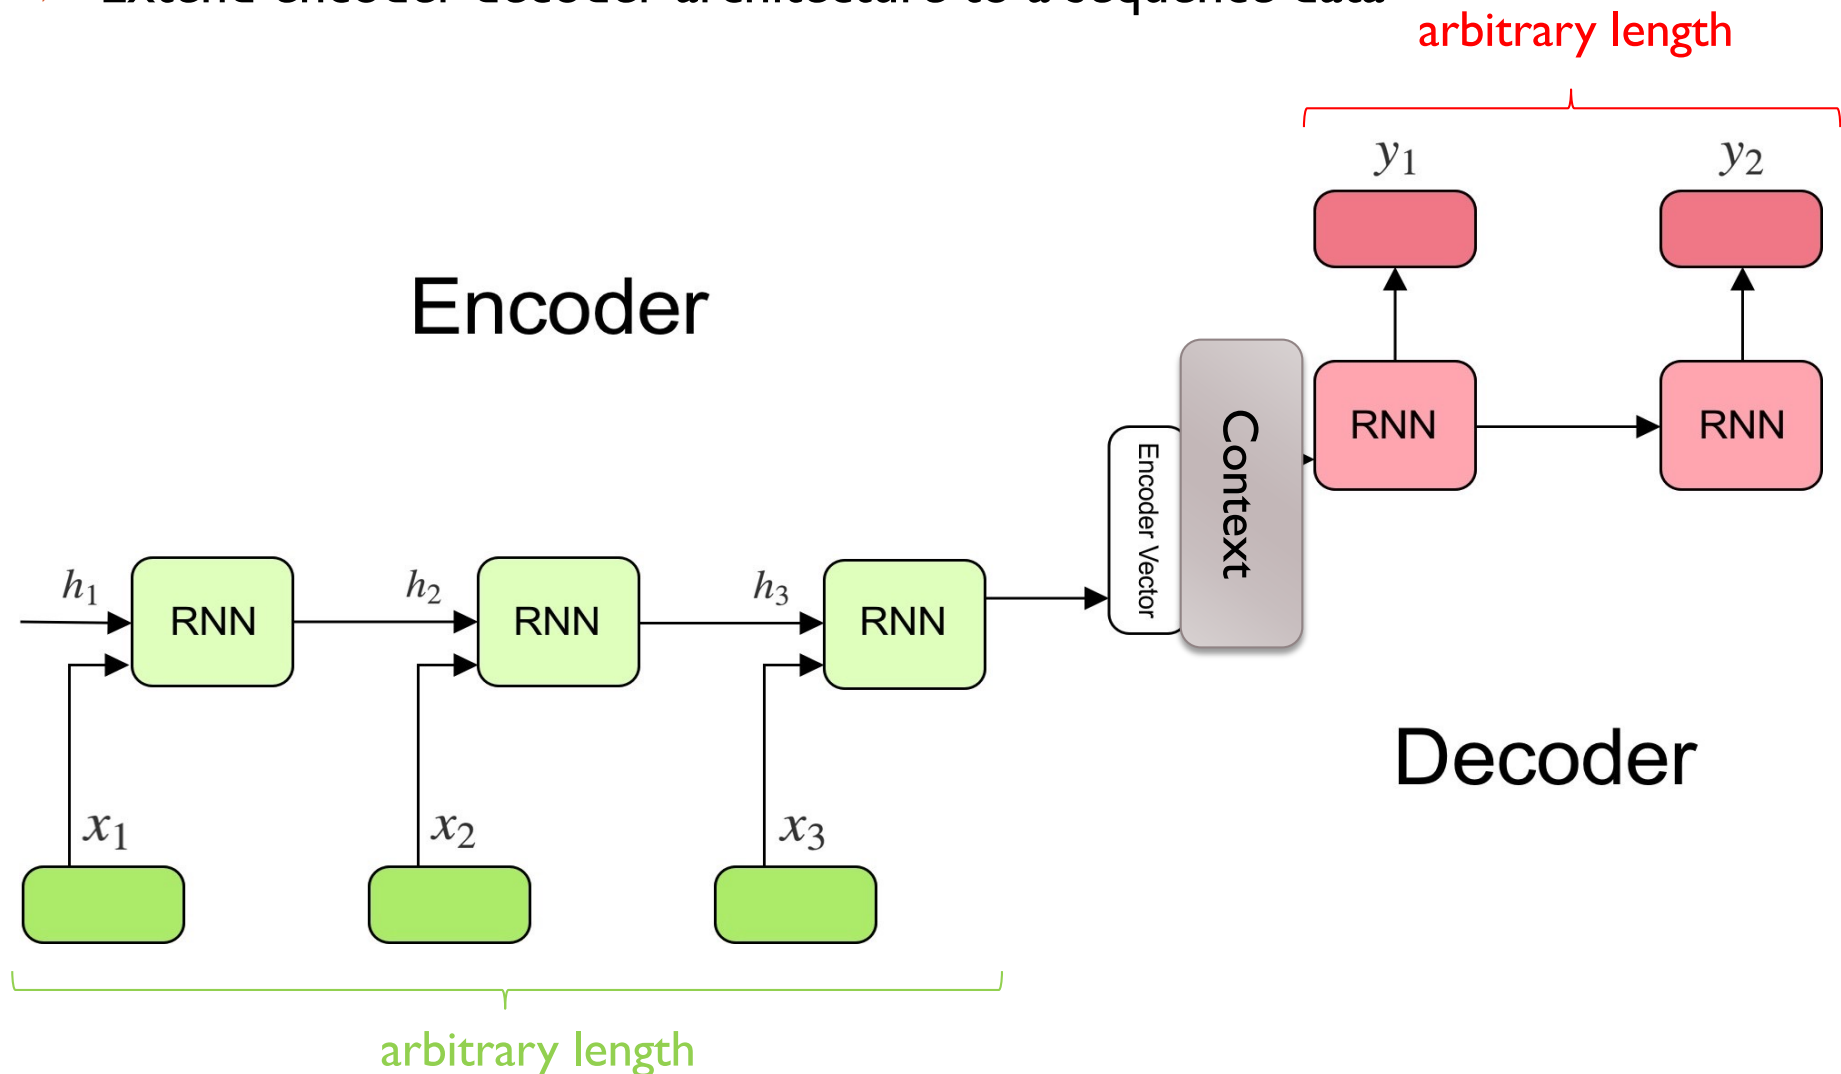


# Image captioning model



# Seq2seq – original approach

- ▶ Extend encoder-decoder architecture to a sequence data



# Sequence-to Sequence Architectures

---

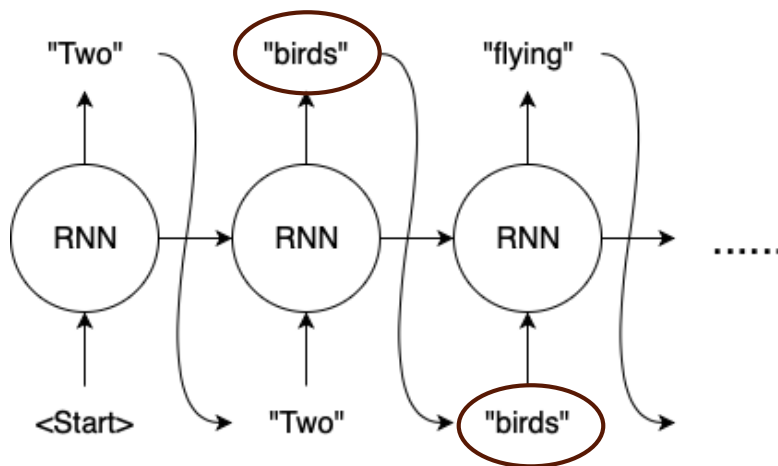
- ▶ Three main part:
  - ▶ **Encoder**: processes the input sequence (ordinary sequence-to-vector RNN)
  - ▶ **Context**: output of the encoder
    - ▶ is usually a simple function of its final hidden state ( $h + c$ )
    - ▶ aims to encapsulate the information for all input elements in order to help the decoder make accurate predictions
  - ▶ **Decoder**: is conditioned on the context to generate the output sequence
    - ▶ the context acts as the initial hidden state of the decoder part of the model
    - ▶ produces output at each step

# How to train seq2seq architecture

## *Without teacher forcing*

`model.fit(input_sent, output_sent)`

Decoder part

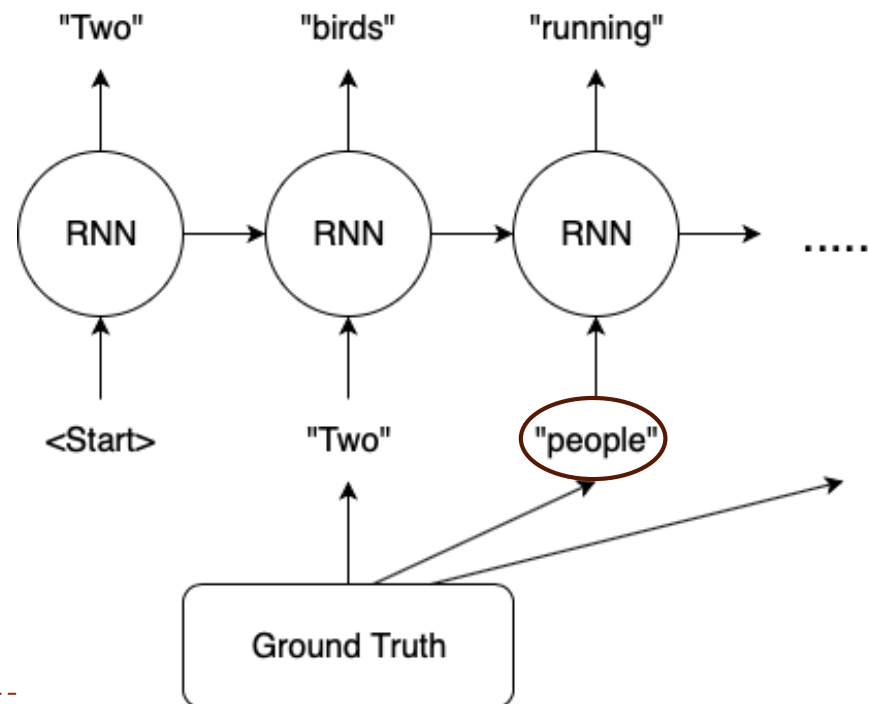


An error is propagated during the training phase → penalizes it

## *With teacher forcing*

`model.fit([input_sent, teacher_sent], output_sent)`

Decoder part

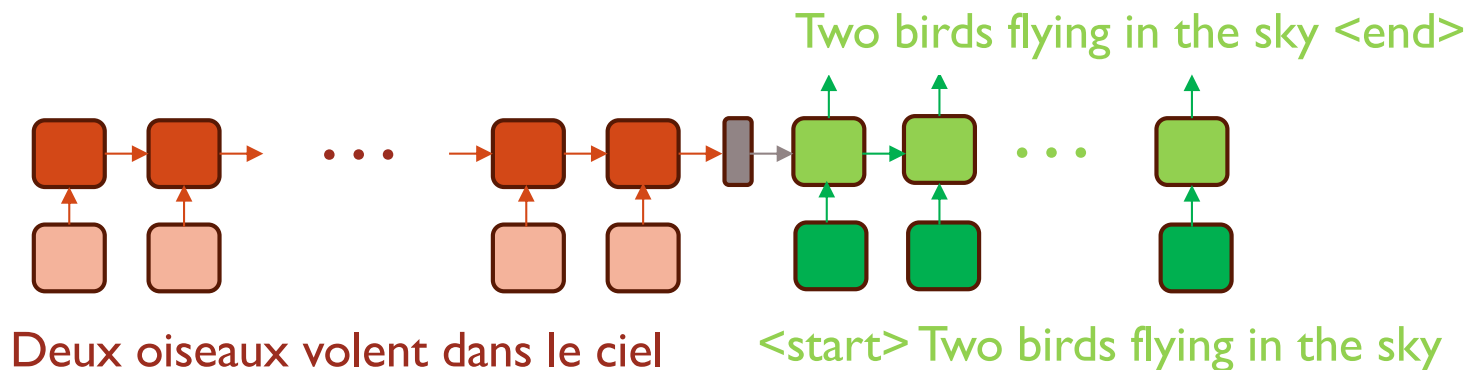


The teacher contains the ground truth



# How to infer with seq2seq architecture

- ▶ Without/ With teacher forcing → no access to a ground Truth
- ▶ Greedy approach
  1. Encode the sentence
  2. iterate to successively decode each time step, reusing the time steps already decoded
    - Then reuse step by step the prediction
- Stop when <end> or max\_length



# Pros and Cons of Teacher Forcing

---

## ▶ Pros:

- ▶ If we do not use Teacher Forcing
  - ▶ the hidden states of the model will be updated by a sequence of wrong predictions
  - ▶ **errors will accumulate**
  - ▶ and it is difficult for the model to learn from that.
- ▶ **Training with Teacher Forcing converges faster.**

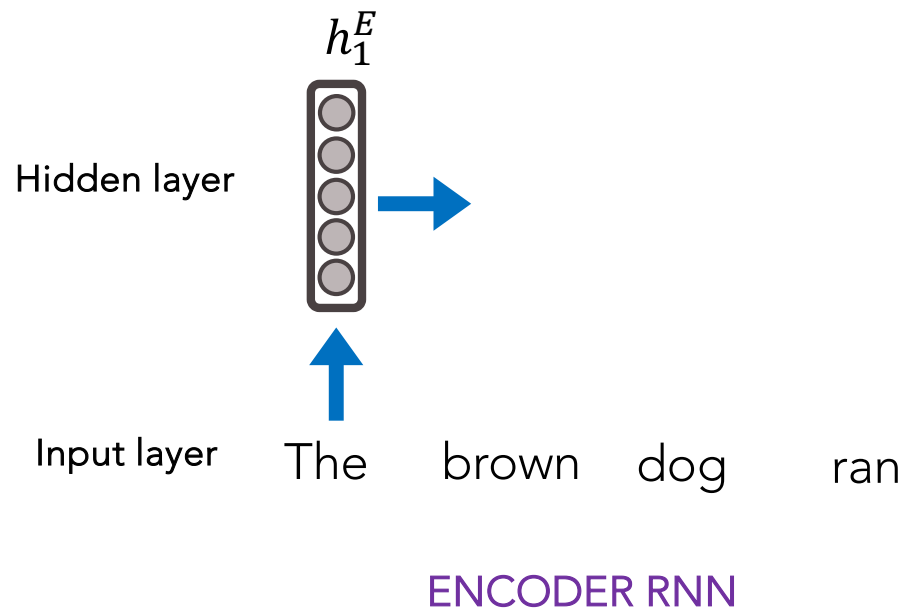
## ▶ Cons:

- ▶ **Unfortunately, during inference, there is no ground truth available**
  - ▶ the RNN model will have to re-inject its own prediction for the next prediction.
- ▶ There is a difference between
  - ▶ learning (no propagation of error)
  - ▶ inference (propagation of error),which leads to poor performance and model instability.
- ▶ This phenomenon is known as **exposure bias** in the literature.



# Train seq2seq model

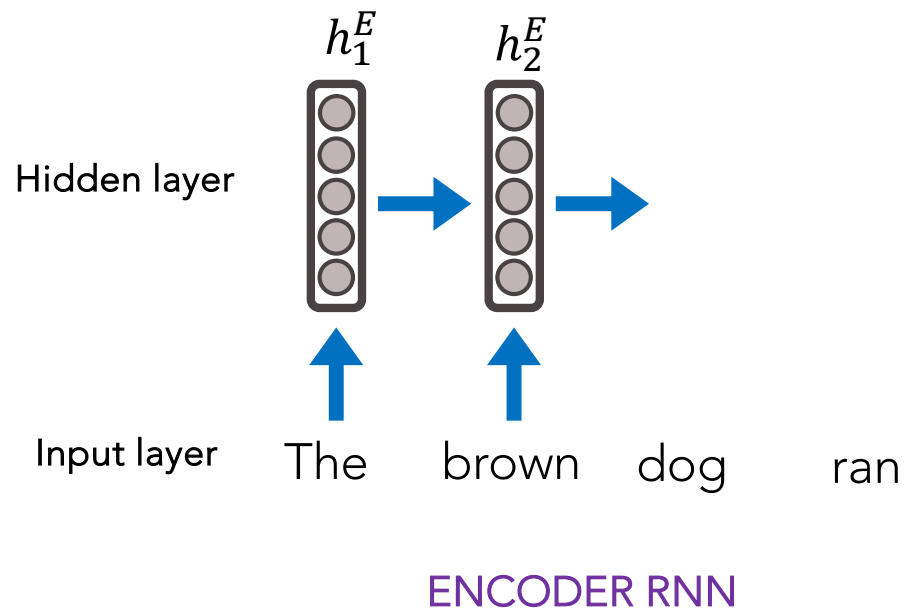
## Step 1 : encode input sentence



# Train seq2seq model

## Step 1 : encode input sentence

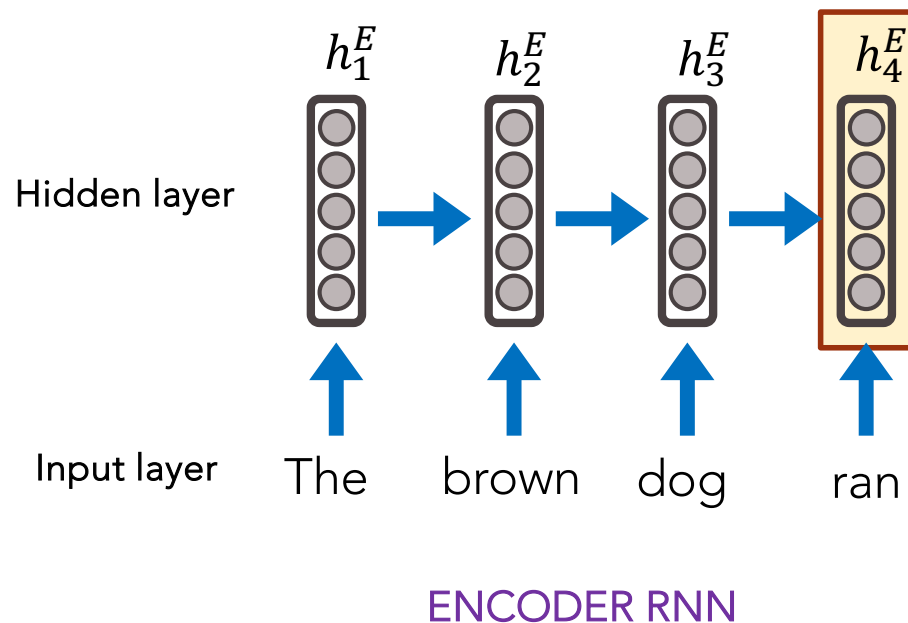
Remember: it's an iterative process until the end of the input sentence



# Train seq2seq model

## Step 1 : encode input sentence

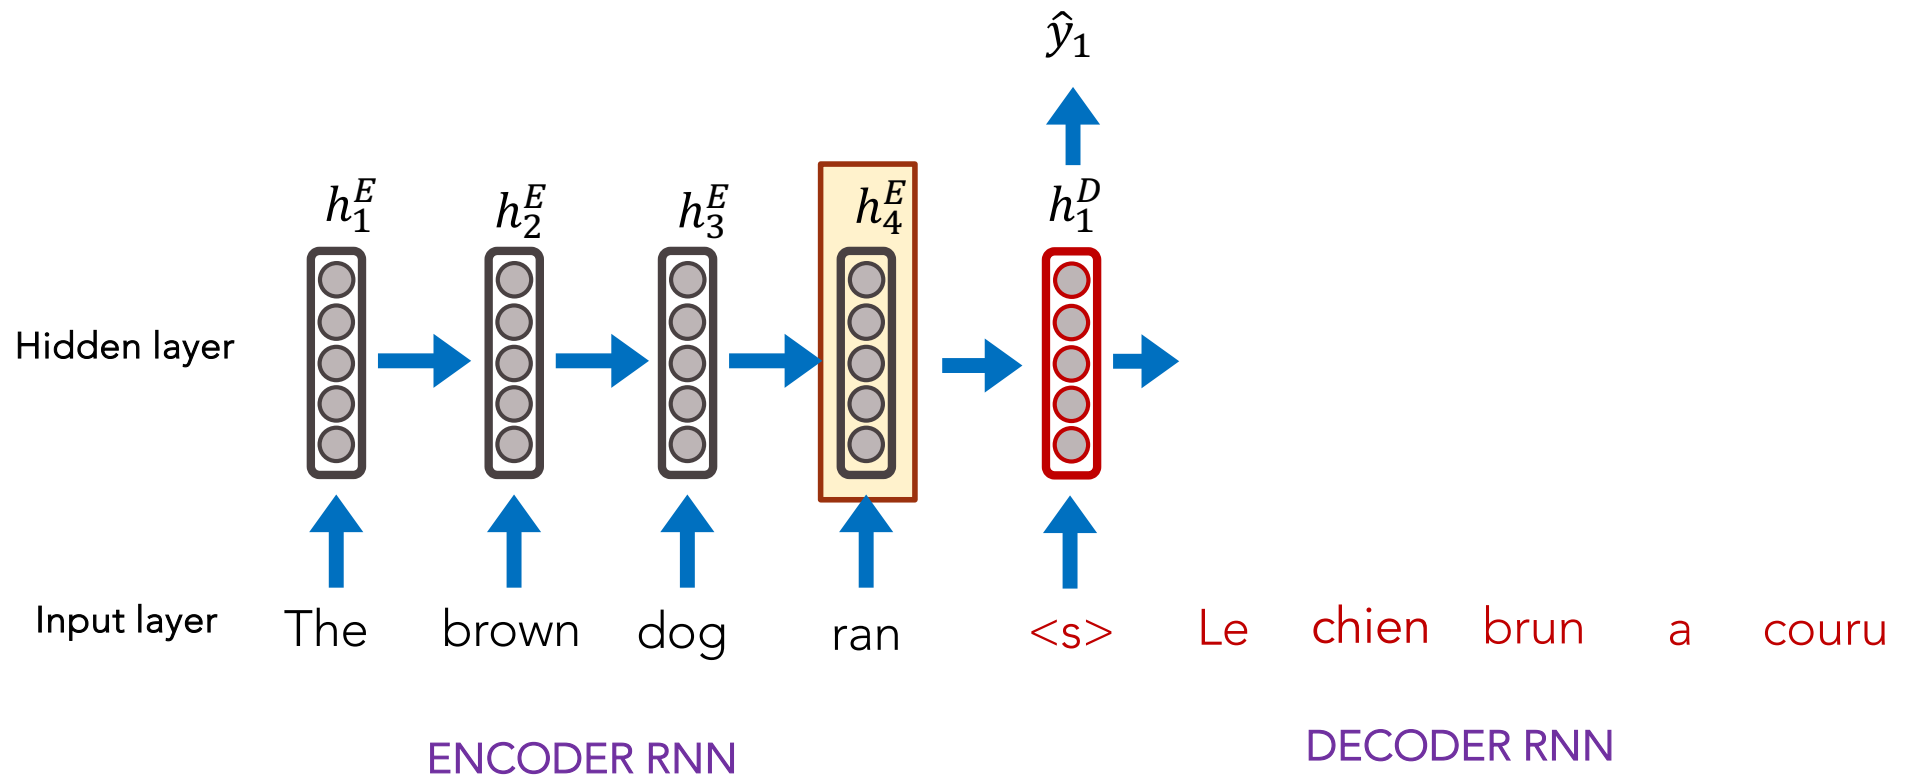
The final hidden state of the encoder RNN  
is the initial state of the decoder RNN



# Train seq2seq model

## Step 2 : decode the sentence

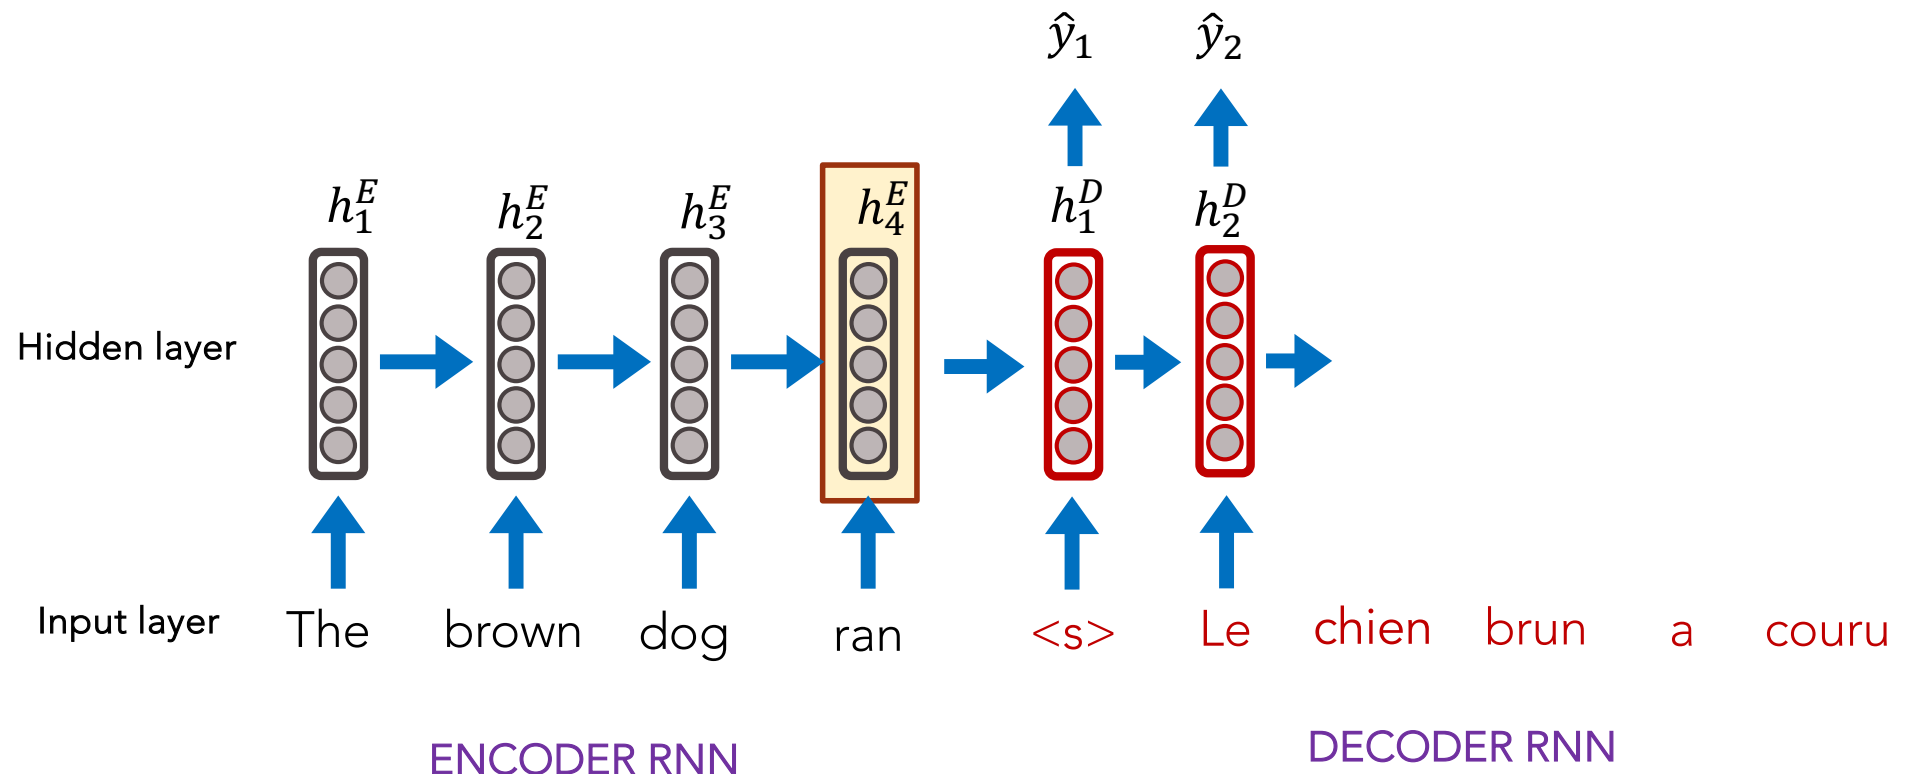
Remember: teacher, help in this task



# Train seq2seq model

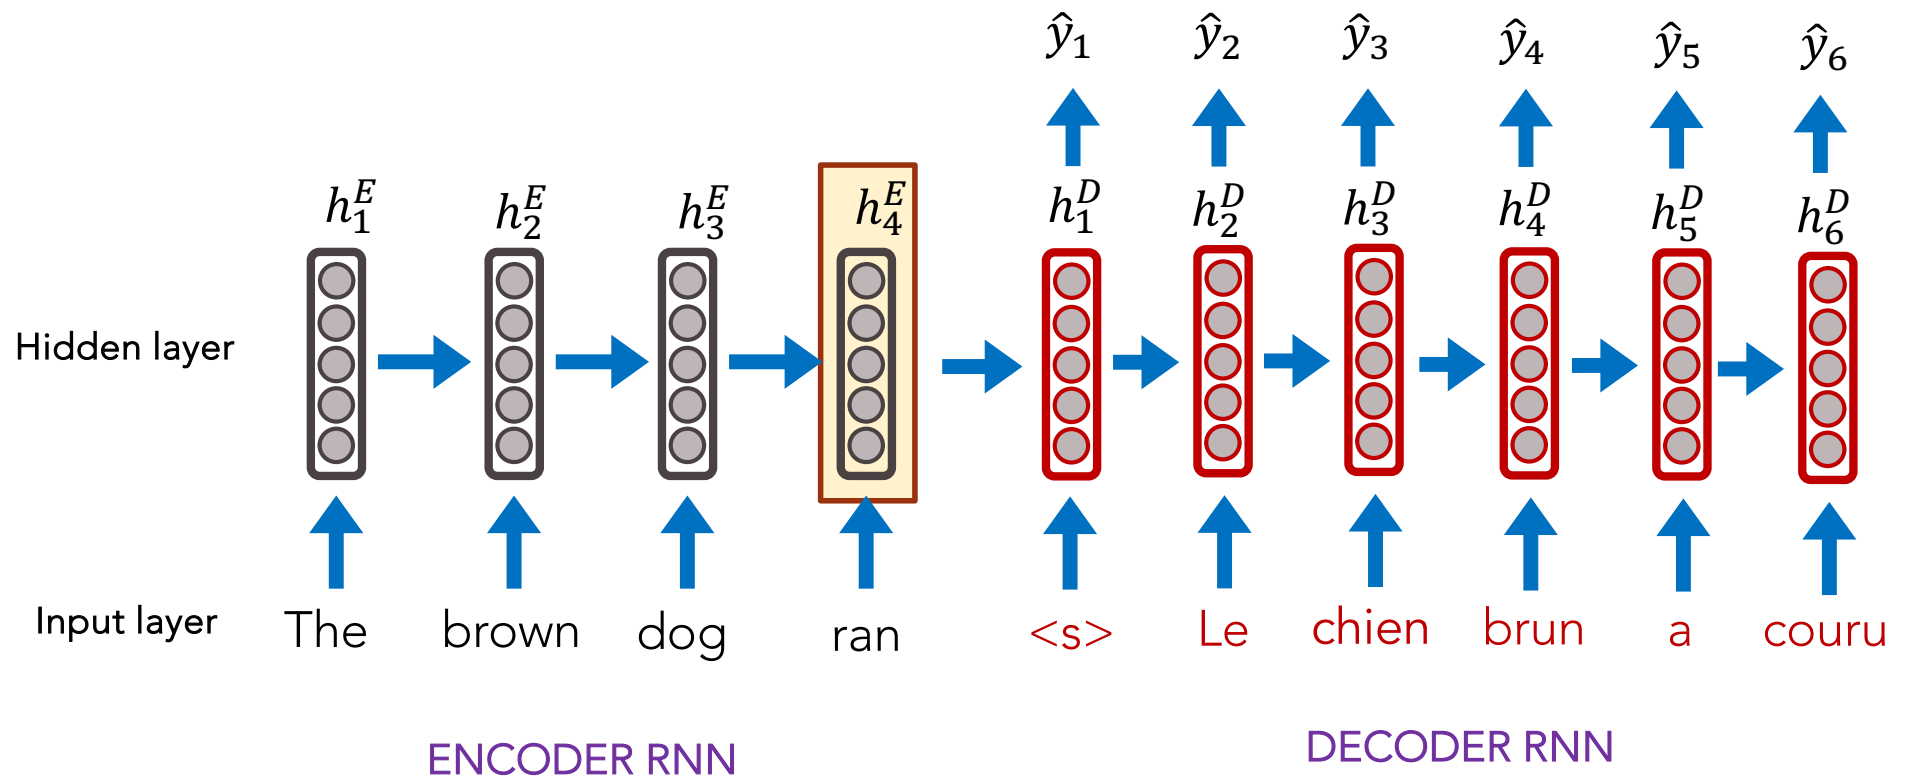
## Step 2 : decode the sentence

Remember: it's also an iterative process  
until the end of the teacher sentence



# Train seq2seq model

## Step 2 : decode the sentence





# How to build Teacher Seq2Seq model

- ▶ For example
  - ▶ Translation from Spanish to English
- ▶ Input sequence = Spanish = (None, None, in\_features)
  - ▶ in\_features: spanish\_vocab\_size=521
  - ▶ Use spanish\_vectorizer
- ▶ Output sequence = English = (None, None, out\_features)
  - ▶ out\_features: english\_vocab\_size = 262
  - ▶ Use english\_vectorizer
- ▶ Embedding dim → as usual (**50**, 100, 150, 300)
- ▶ Latent dim
  - ▶ Represent the size of the latent space (**64**, 128, 256 or more)
  - ▶ Latent space = 2\*latent\_dim for LSTM / latent\_dim for GRU

# How to build Teacher Seq2Seq model

# Define encoder

```
enc_inputs = Input(shape=(None,), name="spanish_input")
```

```
# Why input_shape=(None, None) ?
```

```
enc = Embedding(sp_vocab_size, emb_dim, name="sp_embedding")(enc_inputs)
```

```
_, enc_state_h, enc_state_c = LSTM(latent_dim, return_sequences=False,  
                                   return_state=True)(enc)
```

# Define context

```
context = [enc_state_h, enc_state_c]
```

# How to build Teacher Seq2Seq model

# Define decoder layers

```
layer_embedding = Embedding(en_vocab_size, emb_dim, name="en_embedding")  
layer_lstm = LSTM(latent_dim, return_sequences=True, return_state=True)  
# We use return states in inference.  
layer_dense = Dense(en_vocab_size, activation='softmax')
```

# Define decoder

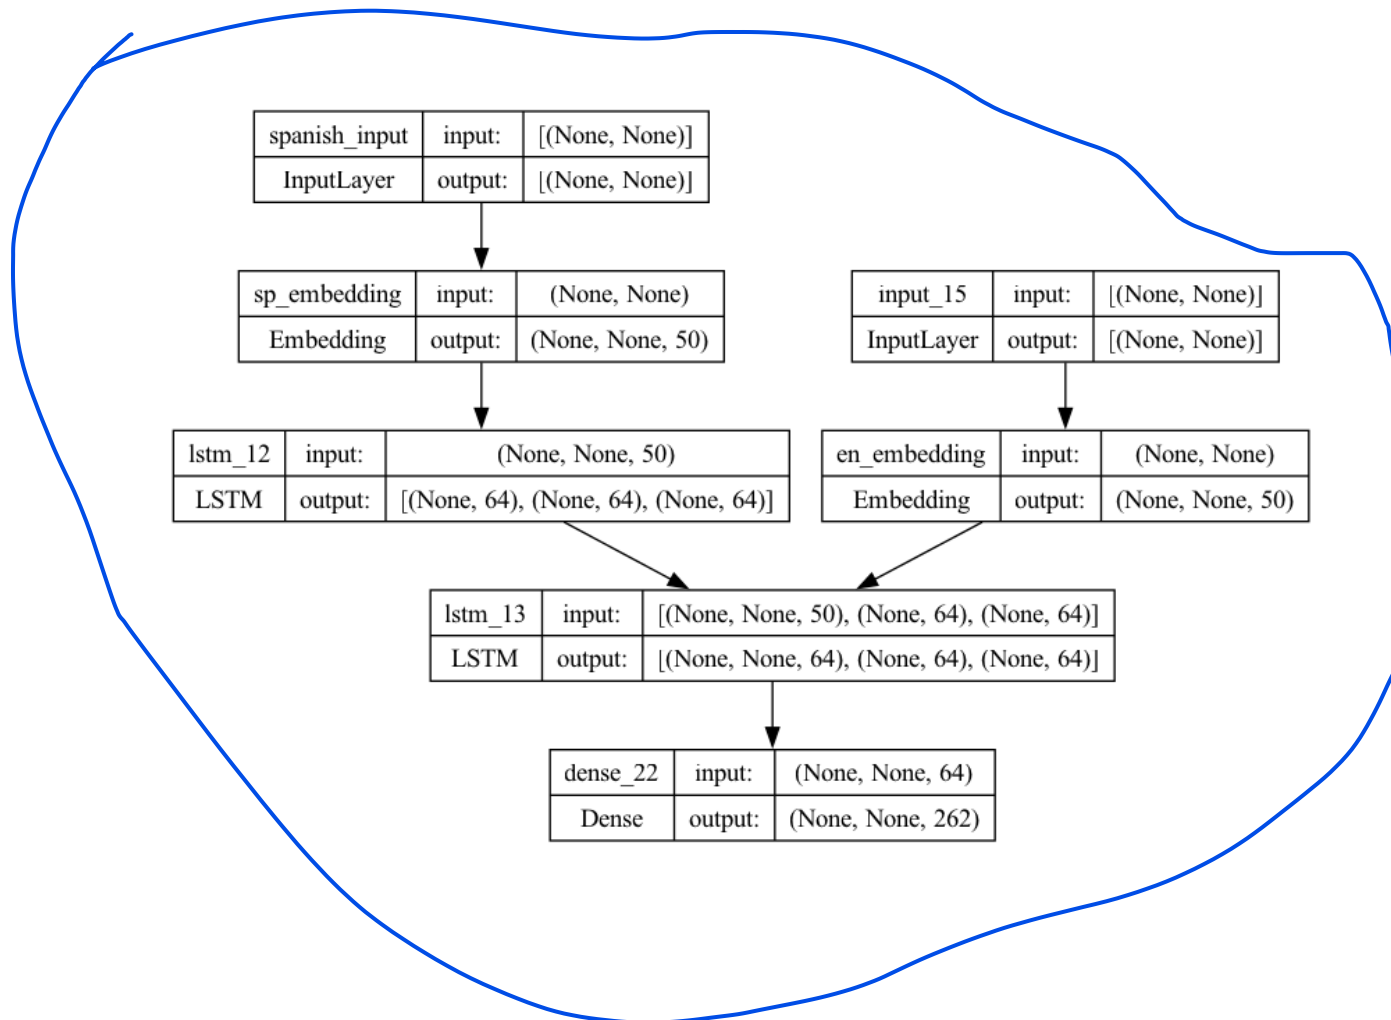
```
dec_inputs = Input (shape=(None,))  
# Why input_shape=(None, None) ?  
dec = layer_embedding (dec_inputs)  
dec, _, _ = layer_lstm(dec, initial_state=context)  
dec_outputs = layer_dense(dec)
```

# Define the Encoder\_Decoder model

```
model = Model ([enc_inputs, dec_inputs], dec_outputs)
```

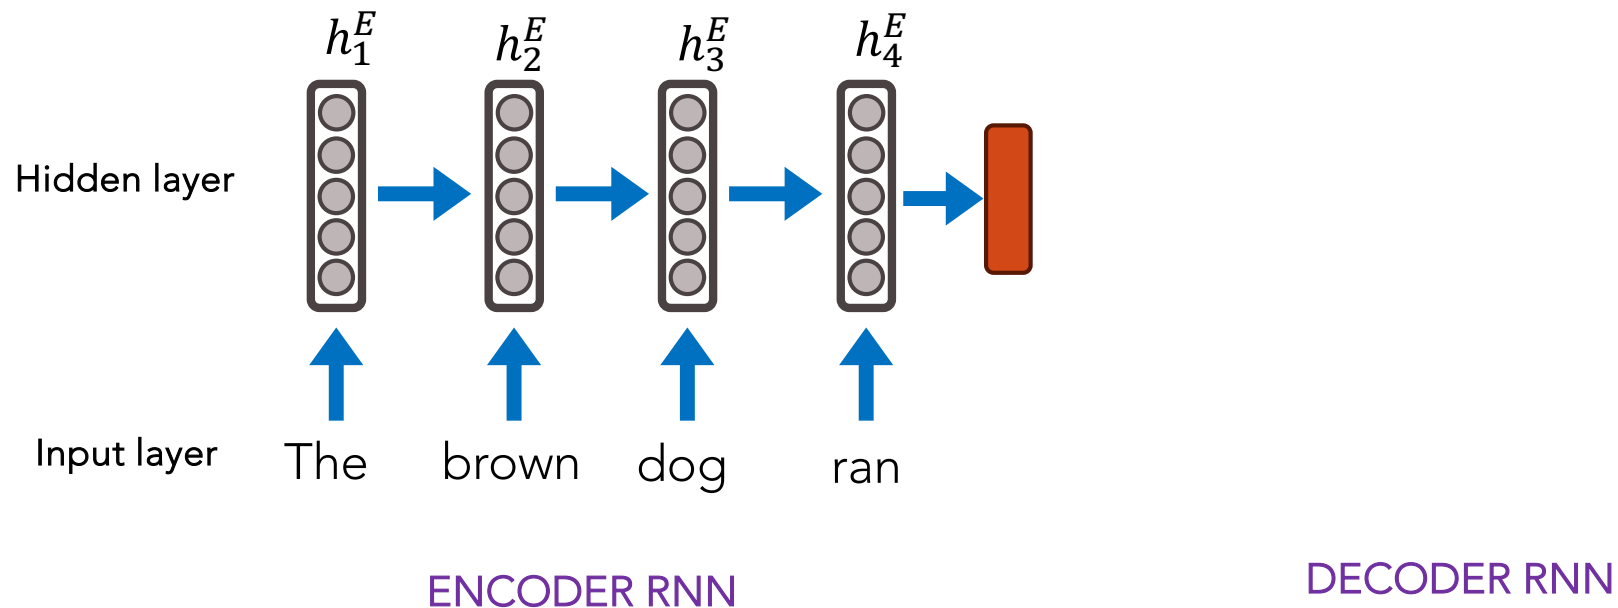


# Seq2Seq model with teacher



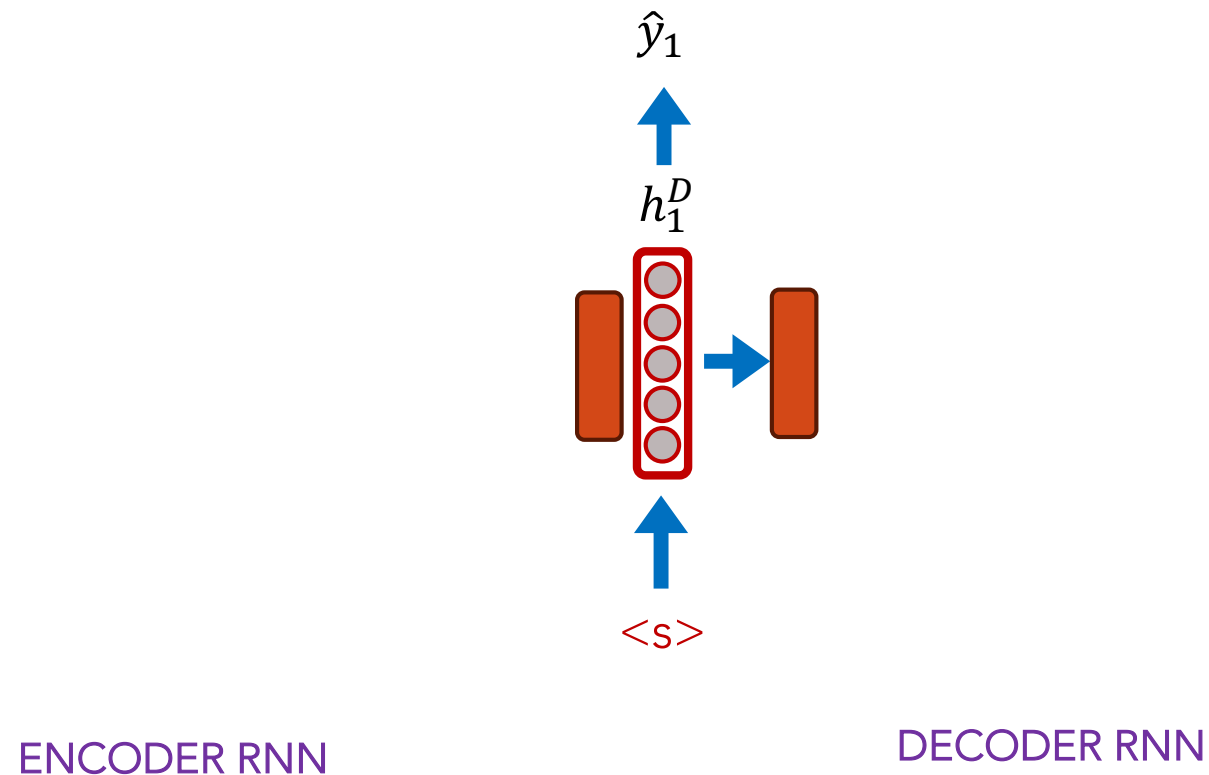
# Predict with seq2seq model

## Step 1 : Use encode to define context



# Predict with seq2seq model

## Step 2 : decode the first output




# Predict with seq2seq model

## Step 2 : step by step... decode sentence

Decode step by step

Reuse at each step, the previous output




# Predict with seq2seq model

## Step 2 : step by step... decode sentence

Decode step by step

Reuse at each step, the previous output

Stop, when generate "stop" label





# How to predict ?

# Build encoder and decoder model

```
encoder_model = Model(enc_inputs, context)
```

```
dec_inputs = Input(shape=(None,))
```

```
dec_input_h = Input(shape=(latent_dim,))
```

```
dec_input_c = Input(shape=(latent_dim,))
```

```
dec = layer_embedding(dec_inputs) # Same cell as previously
```

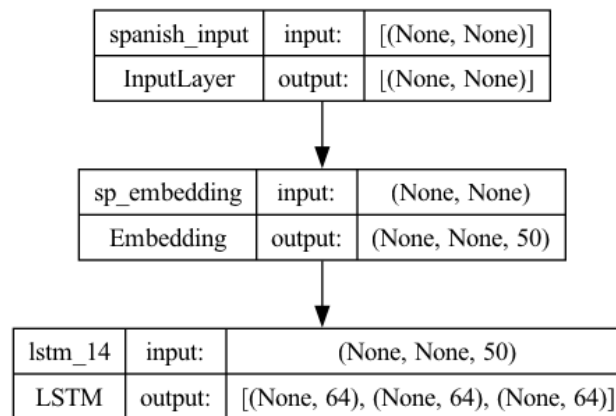
```
dec, dec_h, dec_c = layer_lstm(dec_inputs, # Same cell as  
previously
```

```
initial_state=[dec_input_h, dec_input_c])
```

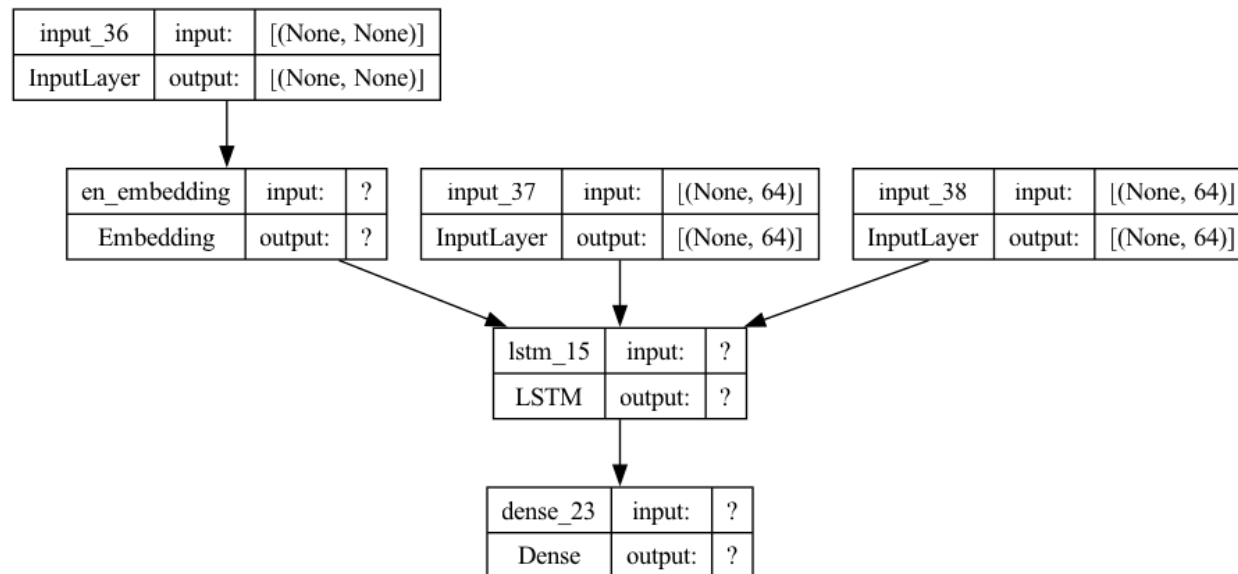
```
decoder_outputs = layer_dense(dec) # Same cell as previously
```

```
decoder_model = Model( [dec_inputs, dec_input_h, dec_input_c],  
[decoder_outputs, dec_h, dec_c])
```

# Encoder model



# Decoder model



# How to predict ?


```
def decode_sequence(input_seq):  
    # Encode the input as state vectors.  
    states_value = encoder_model.predict(input_seq)  
  
    # Iterate over decoded sentence. Target_seq is the input of the decoder  
    target_seq = np.zeros((len(input_seq), 1))  
    target_seq[:, 0] = « initialize the first input »  
  
    output_sequence = [] # Output_sequence is the output of the decoder  
    For _ in range(max_output_length):  
        output_value, h, c = decoder_model.predict( [target_seq] + states_value)  
        # Update the target sequence (of length 1) and state  
        target_seq[:, 0] = decode(output_value)  
        states_value = [h, c]  
        # extend output sequence  
        output_sequence += [target_seq]  
  
    return output_sequence # eventually format it
```

Must necessarily be adapted, as must the architecture of the network according to the problem :

- presence or absence of the embedding layer
- binary classification, categorical classification or mono or multi regression

# Sequence-to-Sequence (seq2seq)

With a Seq2Seq model, we assume that the **entire input sequence** can be **represented by a vector** that is the only interaction between the encoder and the decoder..



# Lab – Build a Deep Learning Translator

---

- ▶ Must necessarily be **finished before next week** as we will continue adding attentions to this model
- ▶ **Dataset:** download your own pair of language and prepare the dataset (code next slide)
  - ▶ <https://www.manythings.org/anki/>
  - ▶ For good performance, it is necessary to have a large dataset
  - ▶ But unfortunately, training a recurrent network is time consuming
  - ▶ We will therefore work with a reduced number of sentences
- ▶ Build a seq2seq neural network
  - ▶ 2 possibilities
    - ▶ At character level
    - ▶ At word level (preferable)
  - ▶ Over-fit your network (very low error rate) with a teacher
    - ▶ We use only a training set (we predict on test)
    - ▶ A very small validation split in order to visualize the overfitting
    - ▶ No EarlyStopping
- ▶ Build model for inference and predict

# Data preparation

---

```
def step1(sent): # sent = on sentence in a language
    def unicode_to_ascii(s): # In order to reduce the possibility
        return ''.join(c for c in unicodedata.normalize('NFD', s) if
unicodedata.category(c) != 'Mn')
    sent = unicode_to_ascii(sent.lower().strip()) # Only lower charater

    # replacing everything with space except (a-z, A-Z, ".", "?", "!", ",",
";")
    sent = re.sub(r"[^a-zA-Z?!.;,]+", " ", sent) # To be adapted
according to the languages chosen

    # creating a space between a word and the punctuation following it.
E.g. "he is a boy." => "he is a boy ."
    # Reference:- https://stackoverflow.com/questions/3645931/python-padding-punctuation-with-white-spaces-keeping-punctuation
    sent = re.sub(r"([?!.;])", r" \1 ", sent) # To be adapted
according to the languages chosen

    return '<start> ' + sent.strip() + ' <end>' # Suppress extra space
```

# Data preparation

---

```
# Loading data
def read_data(path, num_examples):
    # path : path to spa-eng.txt file
    # num_examples : Limit the total number of training example for faster
    training
    lines = io.open(path, encoding='UTF-8').read().strip().split('\n')
    print(lines[0])

    sentences1, sentences2= zip(*[[step1(sent) for sent in
    l.split('\t')[:2]] for l in lines[:num_examples]])

    return np.array(sentences1), np.array(sentences2)
```



# Data preparation

---

```
# Search vocabulary and max_length for each language
def voc(lang):
    # a list of sentences in the same language
    lengths = [len(txt.split()) for txt in lang]
    vocab = set([w for txt in lang for w in txt.split()])

    return max(lengths), list(vocab), len(vocab)+2 # for padding and OOV

max_length1, vocab1, vocab_size1 = voc(sentences1)

# Build vectorizer layer
vectorizer1 = layers.TextVectorization(standardize=None,
                                       output_mode='int',
                                       vocabulary=vocab1,
                                       name= "language1")

# Do the same for language 2
```

