Ensemble methods

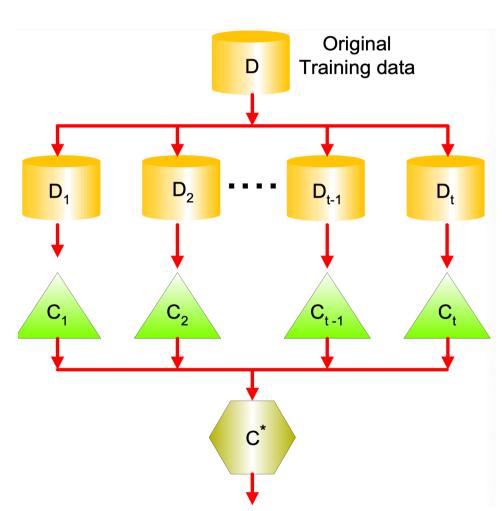
Ensemble Methods

Predict by combining the predictions made by multiple predictors

Step I: build different dataset

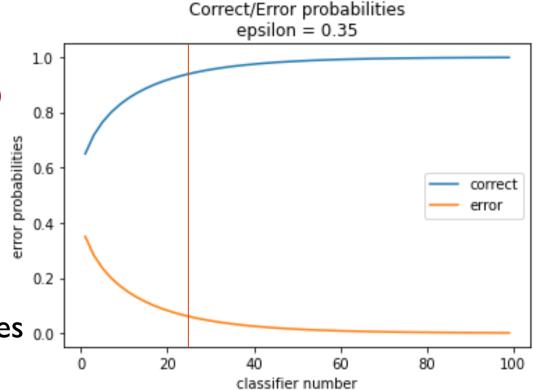
Step2: build different predictors

Step3: combine all predictions



Why Ensemble Methods work?

- Suppose there are 25 base classifiers
 - Each classifier has error rate, $\varepsilon = 0.35$ (<0.5)
 - Assume errors made
 by classifiers are
 uncorrelated
 - Vote for the result
- Probability that the ensemble classifier makes 0.0 a wrong prediction:



$$P(X \ge 13) = \sum_{i=13}^{25} {25 \choose i} \varepsilon^{i} (1 - \varepsilon)^{25 - i} = 0.06$$

Or 94 % gives a correct prediction with a correct rate of 0.65

Ensemble methods

- Useful for classification or regression
 - For classification, aggregate predictions by voting.
 - For regression, aggregate predictions by averaging.
- Model types can be:
 - Heterogeneous
 - ▶ Example: neural net combined with SVM combined with decision tree combined with ...
 - Homogeneous most common in practice
 - ▶ Individual models referred to as base classifiers (or regressors)
 - **Example:** ensemble of 1000 decision trees
- Base classifiers: important properties
 - ▶ Computationally fast: usually need to compute large numbers of classifiers
 - Accuracy: error rate of each base classifier better than random
 - Diversity: lack of correlation
 - \rightarrow independent model \rightarrow they are not wrong on the same examples.

Base classifiers: important properties

Diversity

- Predictions vary significantly between classifiers
- Usually attained by using unstable classifier
 - small change in training data (or initial model weights) produces large change in model structure

Examples of unstable classifiers:

- decision trees (very sensible of dataset)
- neural nets (very sensible of weights initialization)
- rule-based (very sensible of dataset)
- https://machinelearningmastery.com/different-results-each-time-in-machinelearning/

Examples of stable classifiers:

▶ Linear models: linear regression or logistic Regression

How to create diverse base classifiers

Use random projection of the dataset on a lower-dimentional space

- Resample / subsample training data
 - Sample instances
 - Select disjoint partitions
 - □ Reduce the number of training examples
 - Select sample randomly without replacement
 - □ Reduce the number of training examples
 - Select sample randomly with replacement
 - □ Constant number of training examples
 - Sample features (random subspace approach)
 - Select randomly the features used for training

Elementary ensemble algorithm

- S: learning data set of m elements
- 1. Learn the classifier h1 on subset S1 \subset S. Test of h1 on S\S1.
- 2. Learn the classifier h2 on subset S2 \subset S \ S1 with half of elements of S2 wrongly classified by h1.
- 3. Repeat operation until no more remaining element
- 4. H: Majority vote between answers of h1, h2 and h3.
 - Notebook: ensemble voting from scratch.ipynb
- 3. Alternative for 3. Learn classifier h3 on subset S3 \subset S \ S1 \ S2 : contains elements for which rules h1 and h2 answer differently.
- 4. H: Majority vote between answers of h1, h2 and h3.

How to use ensemble in sklearn

- <u>sklearn.ensemble</u>.VotingRegressor
- <u>sklearn.ensemble</u>.VotingClassifier
 - Estimators: list of (str, estimator)
 - Voting: {'hard', 'soft'} # Only for classifier
 - Weights: array-like of shape (n_classifiers,)
- Usual use
 - clf.fit
 - clf.predict
- Use with fitted estimator

```
clf_list = [clf1, clf2, clf3]
eclf = VotingClassifier(estimators = [('1', clf1), ('2', clf2), ('3', clf3)], voting='soft')
eclf.estimators_ = clf_list
eclf.le_ = LabelEncoder().fit(y)
eclf.classes_ = seclf.le_.classes_
# Now it will work without calling fit
eclf.predict(X,y)
```

Main approach of Ensemble Methods

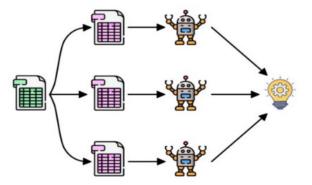
I. Bagging

- Create different training subset from sample training data with replacement
- final output is based on majority voting.

> For example:

> Random Forest

Bagging



➤ 2. Boosting

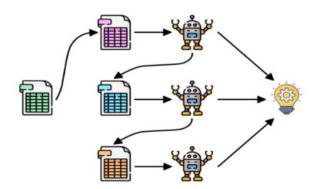
Combines weak learners into strong learners by creating sequential models such that the final model has the highest accuracy.

> For example:

> ADA BOOST: any classifier

> XG BOOST: with tree

Boosting

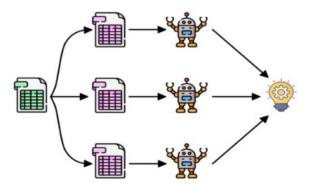


Bagging versus Boosting

Bagging

- Weight of each classifier is same
- Only reduces variance
- Robust to noise and outliers
- Easily parallelized

Bagging

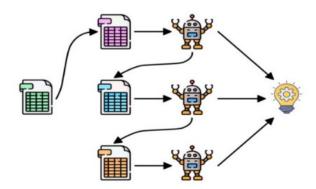


Today: with RandomForest

Boosting

- Weight of a classifier based on its performance
- Reduces both bias and variance
- Sensible to noise and outliers
- Sequential

Boosting



Random Forest an example of bagging method

Original presentation from Jeff Howbert

From decision tree to random forest

- Decision trees are greedy
 - They choose which variable to split on using a greedy algorithm that minimizes error
 - Sensitive to small changes in the data set
 - Overfitting, it's easy
- Combining predictions from multiple trees should work better.
- The random forest
 - Use a subset of data
 - Built different models
- In order to reduce the correlation between the prediction made by each model

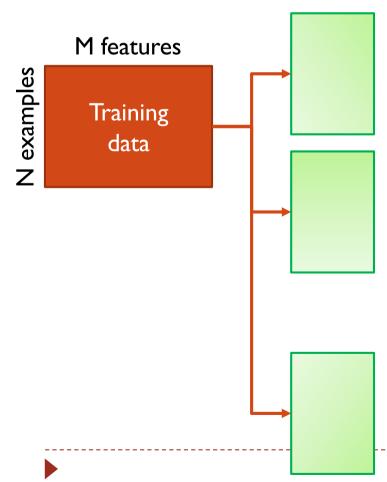
M features Training data

N examples

15

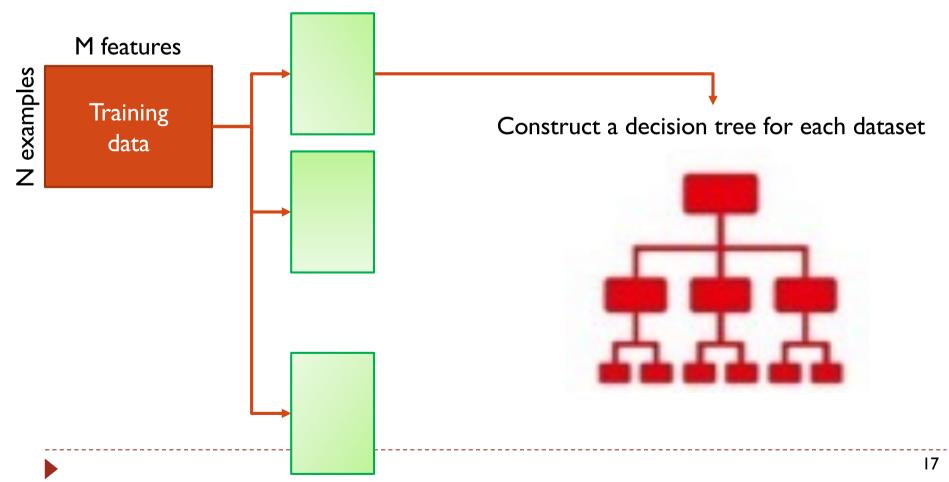
Create many (100) samples from the training data

- with a same number of observations N identical of the original data (remove some samples and duplicate some others bootstrap)
- with m random features (generally $m < \sqrt{M}$)

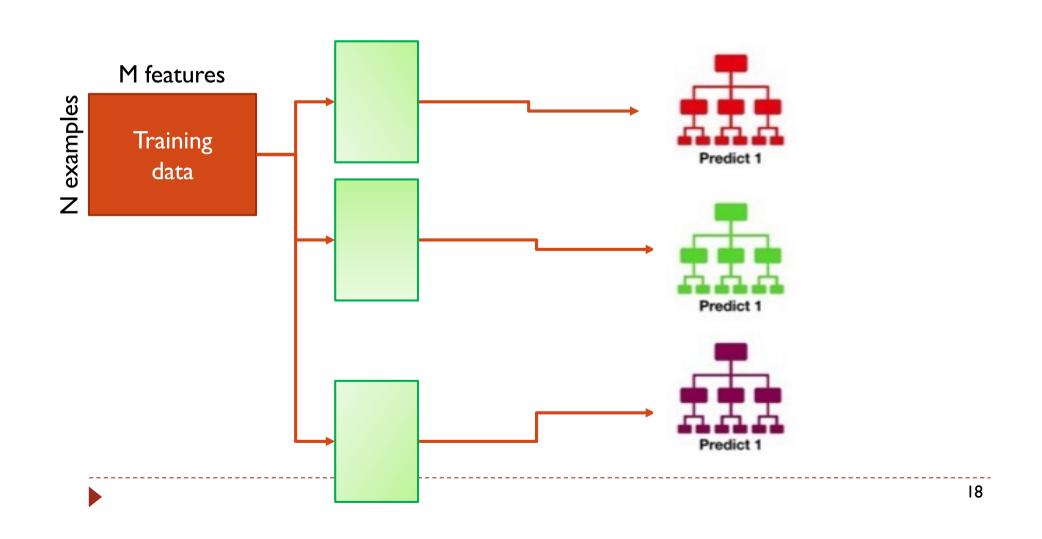


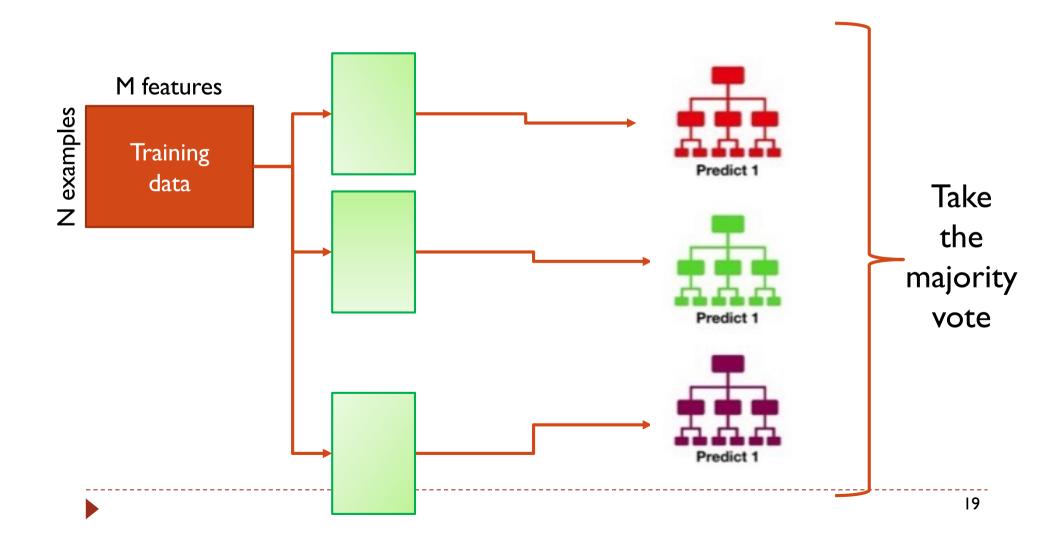
Create many (100) samples from the training data

- with a same number of observations N identical of the original data (remove some samples and duplicate some others bootstrap)
- with m random features (generally $m < \sqrt{M}$)



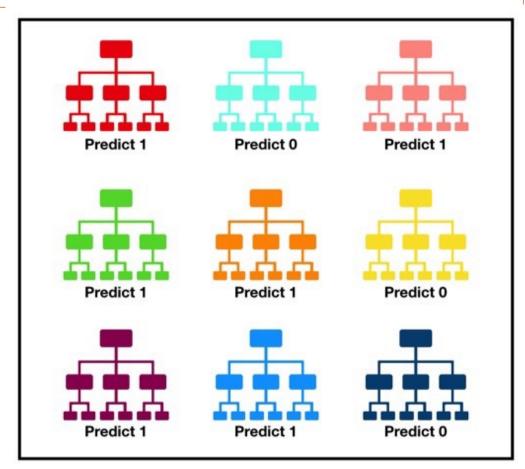
Create decision tree from each bootstrap sample





Why many trees give a more valuable result than one tree

- The reason for this wonderful effect is that the trees protect each other from their individual mistakes.
- While some trees may be wrong, many others will be right.
- Remember
 - With 25 classifier
 - \blacktriangleright Error rate = 0.35
 - ▶ 94 % of correct prediction



Tally: Six 1s and Three 0s

Prediction: 1

Random Forest with sklearn

- Use RandomForestClassifier or RandomForestRegressor
- Main parameters for random forest
 - **n_estimators**: integer, optional (default=10)
 - ▶ The number of trees in the forest.

Random Forest with sklearn

- Parameters of the base model (Decision Tree)
 - max_depth : integer or None, optional (default=None)
 - ▶ The maximum depth of the tree.
 - ▶ If None, then nodes are expanded until all leaves are pure or until all leaves contain less than min_samples_split samples.
 - min_samples_split : int, float, optional (default=2)
 - ▶ The minimum number of samples required to split an internal node
 - min_samples_leaf : int, float, optional (default=1)
 - ▶ The minimum number of samples required to be at a leaf node.
 - A split point at any depth will only be considered if it leaves at least min_samples_leaf training samples in each of the left and right branches.
 - max_features : int, float, string or None, optional (default="auto")
 - ▶ The number of features to consider when looking for the best split
 - min_impurity_decrease : float, optional (default=0.)
 - A node will be split if this split induces a decrease of the impurity greater than or equal to this value.
 - min_impurity_split : float, (default= l e-7)
 - ▶ Threshold for early stopping in tree growth.
 - A node will split if its impurity is above the threshold, otherwise it is a leaf.

PROs and CONs of Random Forest Algorithm

PROs

- I. It solves the problem of overfitting as output is based on majority voting or averaging.
- 2. It performs well even if the data contains null/missing values.
- 3. It is highly stable as the average answers given by a large number of trees are taken.
- 4. It is immune to the curse of dimensionality. Since each tree does not consider all the attributes, feature space is reduced.
- CONs
- I. It is becoming difficult to explain decisions
- 2. long training time if no parallelism

The lab of today

• integrate the use of random forests with a search for hyper-parameters into the project.

Some reading

Some complements to help you revise