



UNIVERSITÉ
CÔTE D'AZUR

Lecture 9: Backpropagation

Optimization for data sciences

Rémy Sun
remy.sun@inria.fr



Course organization

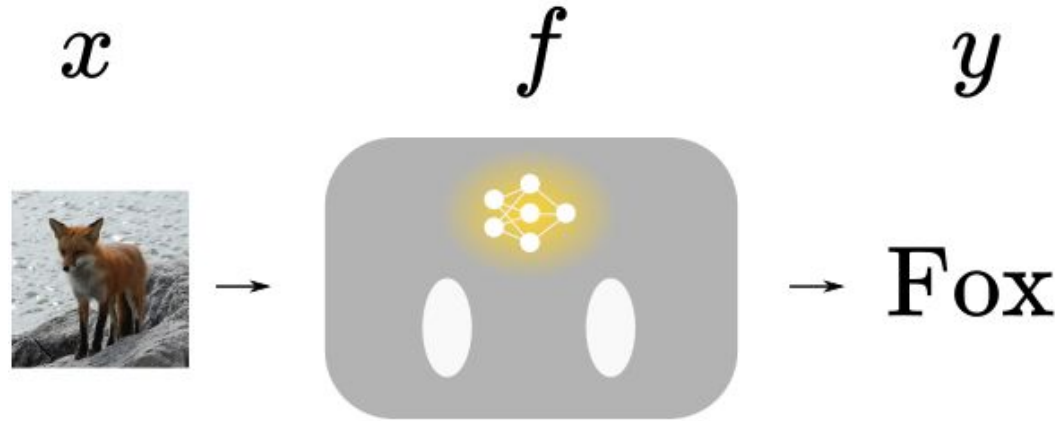
- Introduction to optimization
 - A few problems of interest
 - Quick mathematical refresher
- Convex problems (Following Stephen Boyd)
 - Convex sets
 - Convex functions
 - Convex problems
 - Simplex algorithm for Linear Programming

- Duality (for convex problems)
 - Lagrangian and dual function
 - Dual problem
 - Qualification constraints
 - KKT conditions
- Newton's Descent and Barrier methods for convex case
 - Descent for the unconstrained problems
 - Equality constrained problems
 - Interior point methods
 - Lab session!

- What about the real (neural) world?
 - Problem statement
 - Let's try to solve it!
 - Gradient descent with(out) convexity
 - Gradient descent variants
- **Backpropagation**
 - **Chain rule derivation**
 - **Dynamic programming**
 - **Backpropagation**
 - **Lab session 2!**

- Reports on lab sessions
 - Labs on jupyter notebooks
 - Not every session
 - Explain the code done in the session
 - Summarize what is done in the practical
- Written Exam
 - Theoretical questions
 - We will do exercises in class

Refresher on convexity!

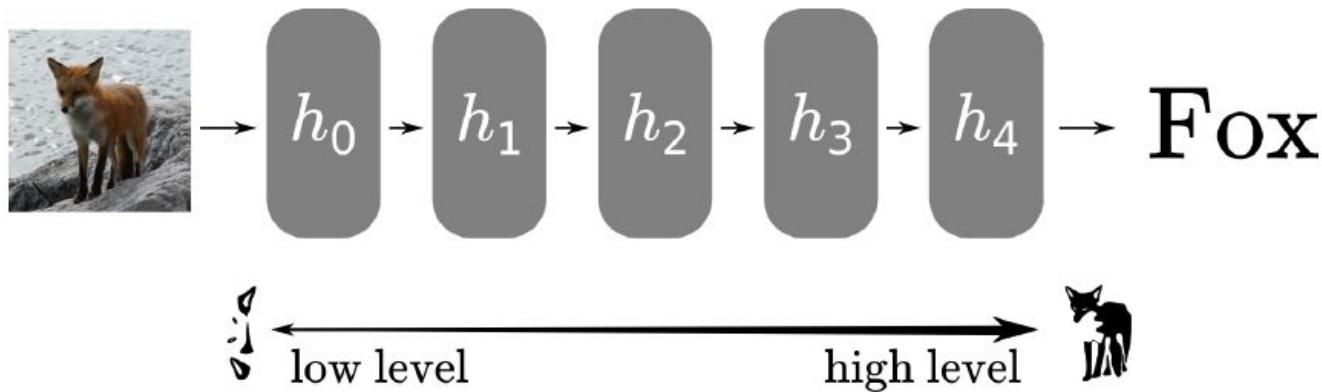


- Find (robot) f that classifies images well
 - Often based on neural networks

$$\forall (x, y) \in \mathcal{D}, f(x) = y$$

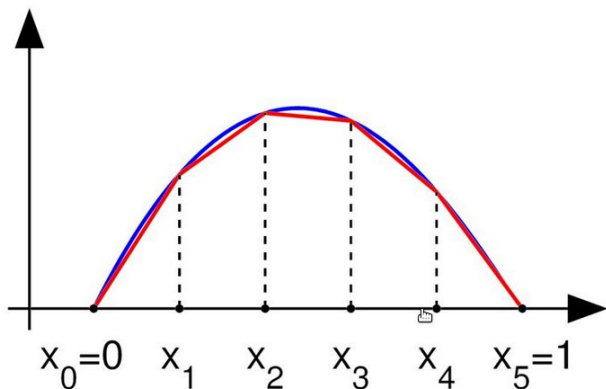

- Problem: we do not know \mathcal{D} !
 - Solved problem otherwise...
 - Evaluating the risk requires this distribution
- Solution: Use a dataset D of (x,y) sampled from \mathcal{D}
 - **Empirical Risk Minimization**
 - If the (x,y) are i.i.d drawn from \mathcal{D} can be expressed as a mean over the dataset

$$\min_{\theta} \hat{\mathcal{R}}_{\theta} = \frac{1}{N} \sum_{i=0, \dots, N-1} l(f_{\theta}(x_i), y_i)$$



- Neural networks are sequences of simple functions

$$f_{\theta} = h_{\theta}^0 \circ h_{\theta}^1 \circ \dots \circ h_{\theta}^{L-1}$$



- Highly expressive
 - Can fit many types of distributions

$$\begin{array}{ll}\text{minimize} & f_0(x) \\ \text{subject to} & f_i(x) \leq 0, \quad i = 1, \dots, m \\ & h_i(x) = 0, \quad i = 1, \dots, p\end{array}$$

- Cost function is given by neural network and loss
- No constraints!
 - Easy unconstrained problem!
- **Not convex, for a number of reasons.**

- What do we need?
 - Step size
 - Fixed step size
 - Gradient
 - Could get pretty expensive too...
 - Obtained through backpropagation!
 - No Hessian!

- Under reasonable L-smooth assumption
 - Gradient descent converges!
 - To something
 - With a minuscule fixed step size
- General deep learning heuristics
 - Most local minimums are similarly good/bad
 - Big networks have very few really bad mins

1. Some context

$$\theta^{t+1} := \theta^t - \eta \nabla_{\theta} \mathcal{R}_{\theta}(\hat{B})$$


- Requires finding the risk gradient wrt parameters

$$\nabla_{\theta} \mathcal{R}_{\theta}(\hat{B}) = \frac{1}{\#B} \sum_{k=0, \dots, B-1} \nabla_{\theta} l(f_{\theta}(x_k), y_k)$$

- Boils down to computing gradients for one sample


$$\nabla_{\theta} l(f_{\theta}(x), y)$$

Let's get the gradients!



- Neural networks embed complex functions
 - Reason we use them


Let's get the gradients!



- Neural networks embed complex functions
 - Reason we use them
- **How do we get the gradients?!**


- Not an easy problem originally
- First neural networks do not use gradients
 - “Hard” value perceptron
 - Rosenblatt’s algorithm
- 1970: Modern Backpropagation (reverse mode)
 - Computes gradients
 - Uses Chain rule derivation
 - With bottom-up dynamic programming

2. Chain rule



- The complete network is complicated

$$f_{\theta} = h_{\theta}^0 \circ h_{\theta}^1 \circ \dots \circ h_{\theta}^{L-1}$$



- But individual derivatives are simple!

$$f_{\theta} = h_{\theta}^0 \circ h_{\theta}^1 \circ \dots \circ h_{\theta}^{L-1}$$

$$\frac{dz}{dx} = \frac{dz}{dy} \cdot \frac{dy}{dx}$$

- Derivatives can be broken down into simpler parts
 - As a product!
- Allows us to use the very simple functions
 - E.g. $\tilde{h}_i = \sum_{j=1}^{n_x} W_{i,j}^h x_j + b_i^h$ has trivial partial derivatives
- Known for a very long time (17th century)


- Chain rule derivation powers backpropagation
 - Leverages simple component functions
 - Easy to break down
- Works as long as we know how to derivate layers
- Backpropagation is not just the chain rule
 - Exploding costs
 - Lots of redundant computations

- Chain rule derivation powers backpropagation
 - Leverages simple component functions
 - Easy to break down
- Works as long as we know how to derivate layers
- **Backpropagation is not just the chain rule**
 - Exploding costs
 - Lots of redundant computations
 - Backprop is **efficient** derivation

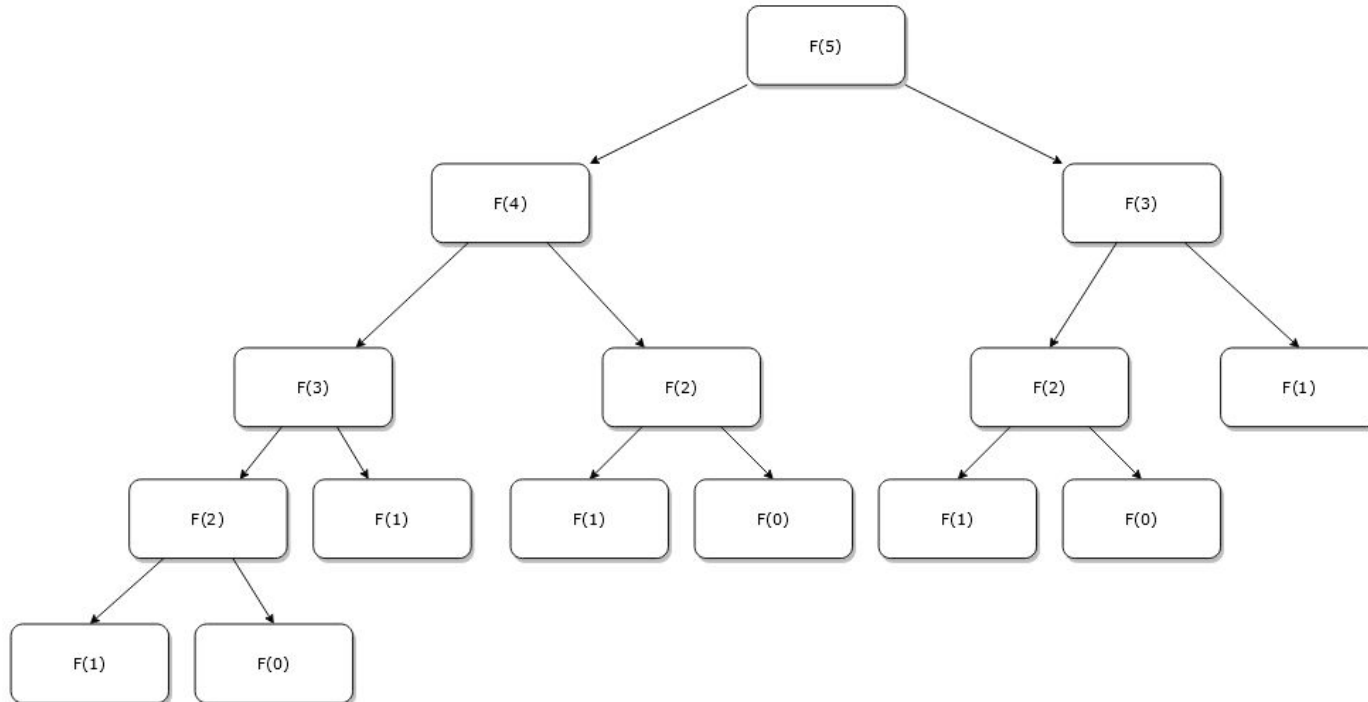
3. Dynamic Programming

- Efficient computations of complex problem
 - Bellman (1950)
 - Break down into simpler (re-occurring) problems
 - And remember your solutions
- Avoid redundant computations
 - By using additional memory to store results
 - By properly structuring computations

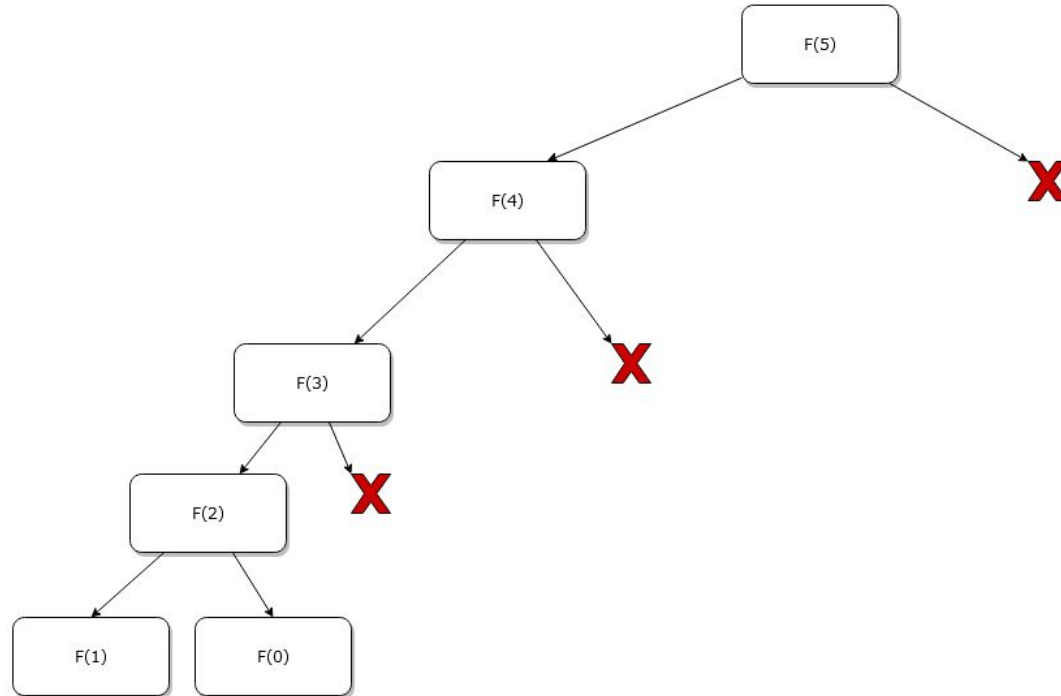
- $F(0)=0$, $F(1)=1$, $F(N+2)=F(N)+F(N+1)$




Lots and lots and lots of computations...



- Memoization: Store sub-problem solutions as needed



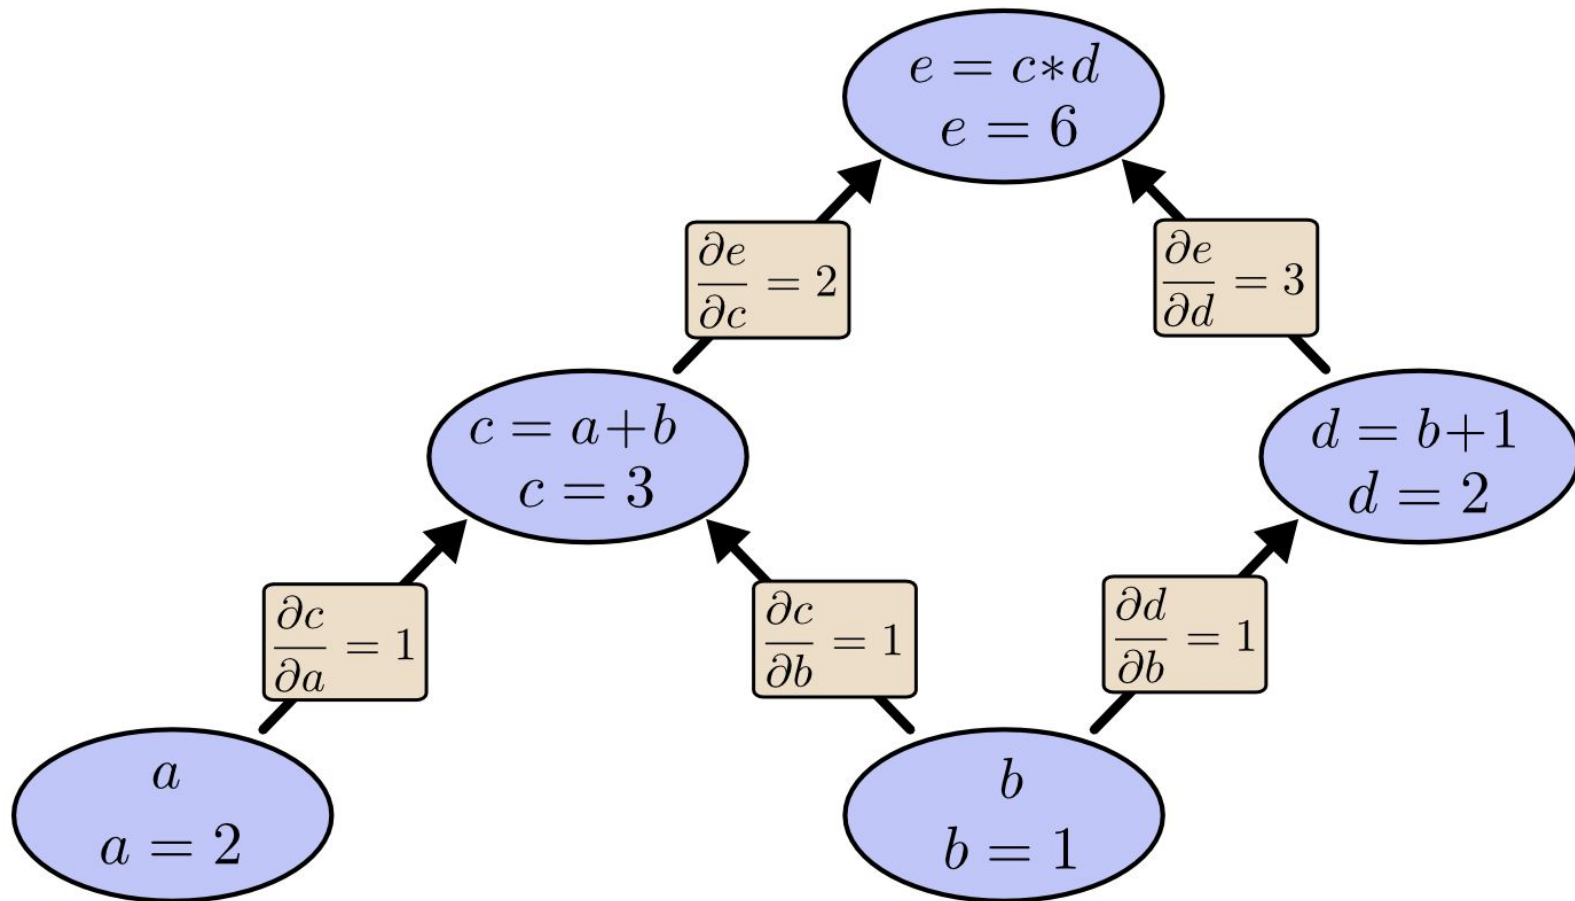
- Tabulation: Start from basic problems
 - Tabulate results as we go along
- One single forward pass this way!
- Memoization vs. Tabulation depends on the problem



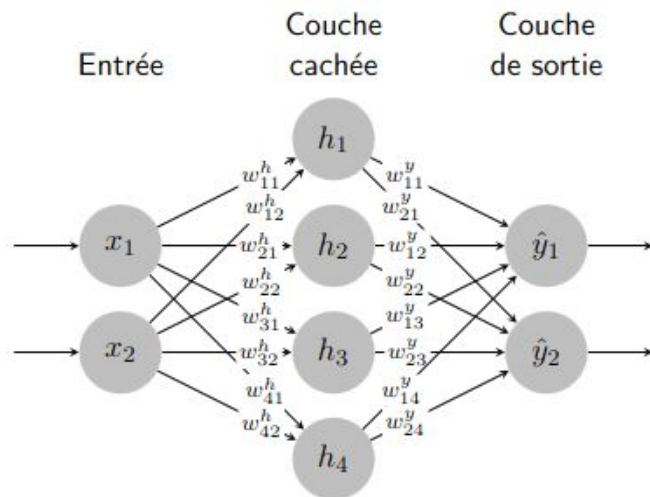
- Efficient computations of complex problem
 - Bellman (1950)
- Avoid redundant computations
 - By using additional memory to store results
 - By properly structuring computations
- Top-down (Memoization): Start from final problem
- Bottom-up (Tabulation): Start from basic problem

4. Backpropagation

- Modern Backpropagation: Linnainmaa (1970)
 - In his Master's thesis!
 - Some precursor efforts before
 - Theorized by Rosenblatt for perceptrons
- Efficient differentiation in a computational graph
 - Reverse mode autodiff
 - Not for neural network per se
 - Applied to neural networks later on



- Networks are complex but made of simple parts!
 - Simple gradients of component functions
 - Chain-rule allows decomposition into simple gradients $\frac{\partial l}{\partial w} = \frac{\partial l}{\partial a} \frac{\partial a}{\partial w}$
 - Remember the computed values
- Need to store intermediate activations “a” to evaluate partial derivatives $\frac{\partial a}{\partial w}$
- Only one pass (in backward)!

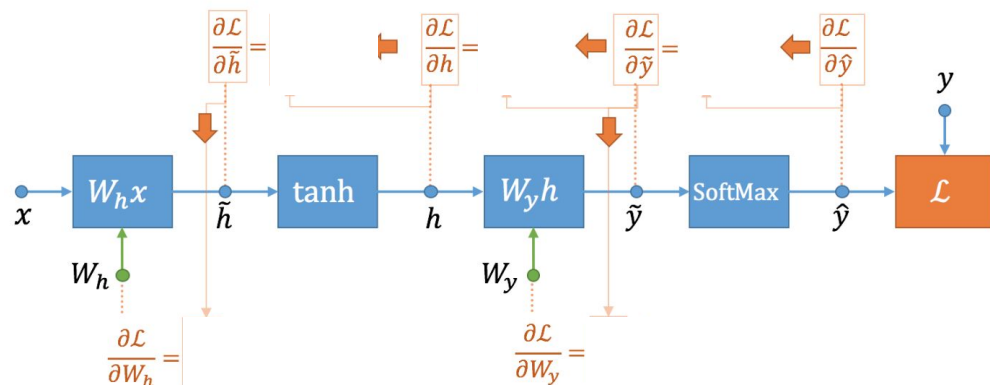
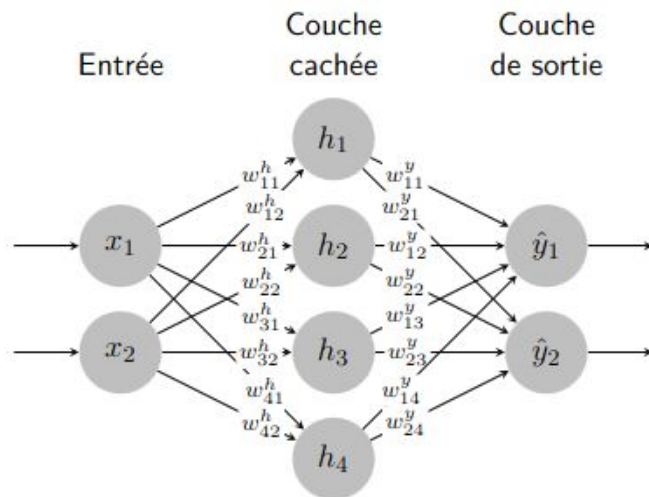


- Simple 1 hidden layer MLP
 - 2 inputs
 - 2 outputs
 - 4 hidden activations
- Classification problem
 - Outputs probabilities
 - Cross-entropy loss

$$l_{CE}(\hat{y}, y) = - \sum_{i=0}^{\#Classes-1} y_i \log(\hat{y}_i)$$

Example: Tanh MLP

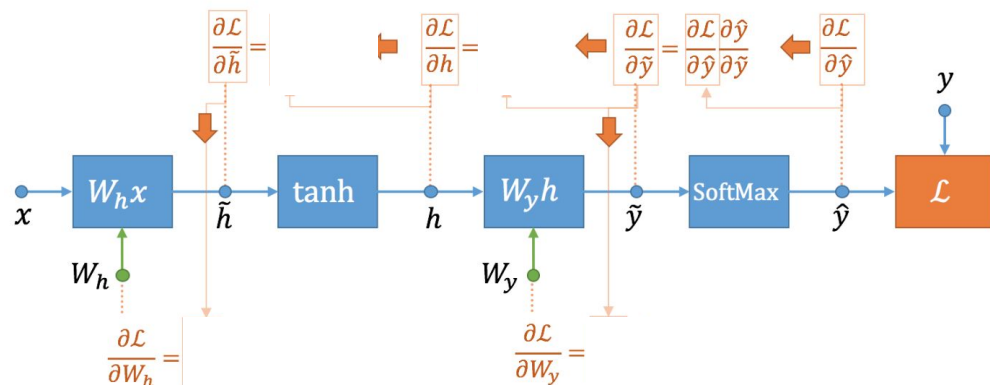
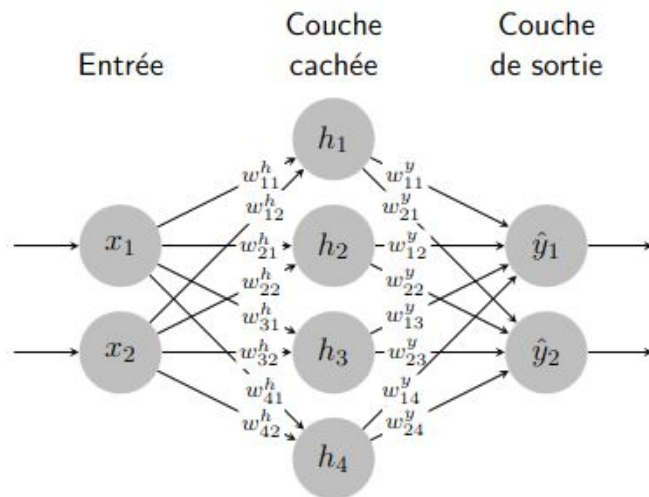
$$l_{CE}(\hat{y}, y) = - \sum_{i=0}^{\#Classes-1} y_i \log(\hat{y}_i)$$



$$\begin{cases} \tilde{h}_i = \sum_{j=1}^{n_x} W_{i,j}^h x_j + b_i^h \\ h_i = \tanh(\tilde{h}_i) \\ \tilde{y}_i = \sum_{j=1}^{n_h} W_{i,j}^y h_j + b_i^y \\ \hat{y}_i = \text{SoftMax}(\tilde{y}_i) = \frac{e^{\tilde{y}_i}}{\sum_{j=1}^{n_y} e^{\tilde{y}_j}} \end{cases}$$

Example: Tanh MLP

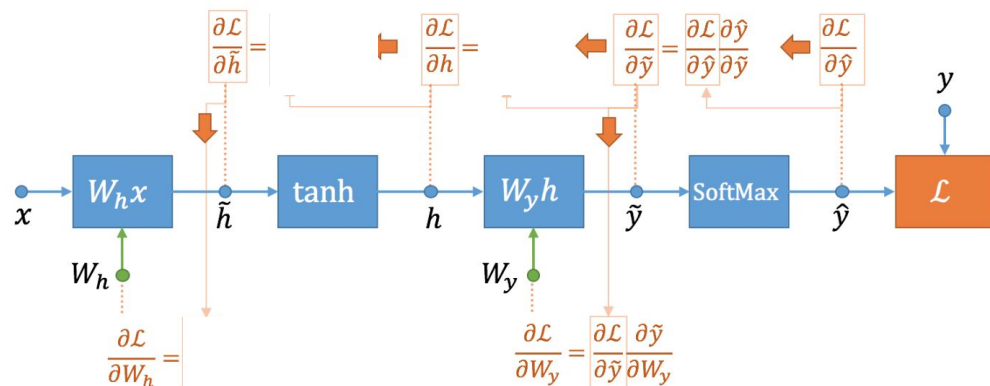
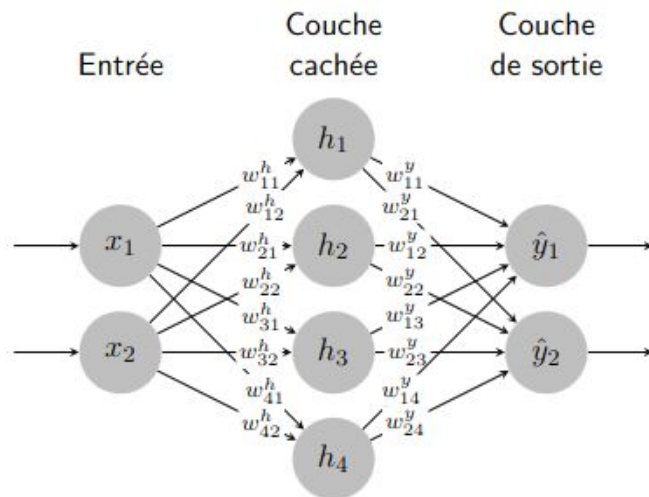
$$l_{CE}(\hat{y}, y) = - \sum_{i=0}^{\#Classes-1} y_i \log(\hat{y}_i)$$



$$\begin{cases} \tilde{h}_i = \sum_{j=1}^{n_x} W_{i,j}^h x_j + b_i^h \\ h_i = \tanh(\tilde{h}_i) \\ \tilde{y}_i = \sum_{j=1}^{n_h} W_{i,j}^y h_j + b_i^y \\ \hat{y}_i = \text{SoftMax}(\tilde{y}_i) = \frac{e^{\tilde{y}_i}}{\sum_{j=1}^{n_y} e^{\tilde{y}_j}} \end{cases}$$

Example: Tanh MLP

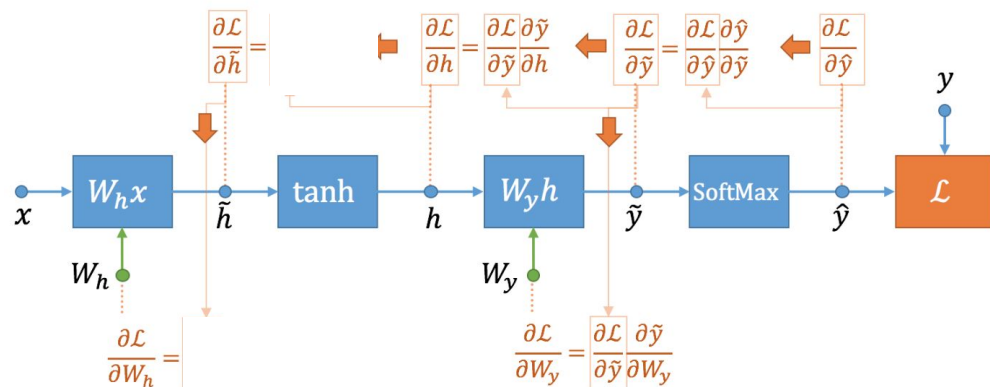
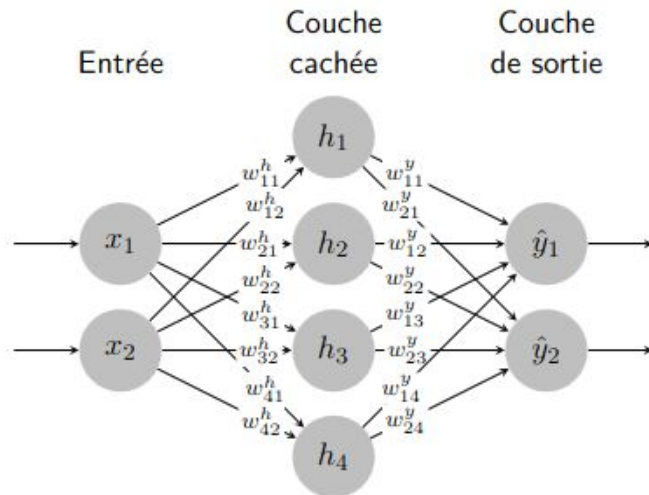
$$l_{CE}(\hat{y}, y) = - \sum_{i=0}^{\#Classes-1} y_i \log(\hat{y}_i)$$



$$\begin{cases} \tilde{h}_i = \sum_{j=1}^{n_x} W_{i,j}^h x_j + b_i^h \\ h_i = \tanh(\tilde{h}_i) \\ \tilde{y}_i = \sum_{j=1}^{n_h} W_{i,j}^y h_j + b_i^y \\ \hat{y}_i = \text{SoftMax}(\tilde{y}_i) = \frac{e^{\tilde{y}_i}}{\sum_{j=1}^{n_y} e^{\tilde{y}_j}} \end{cases}$$

Example: Tanh MLP


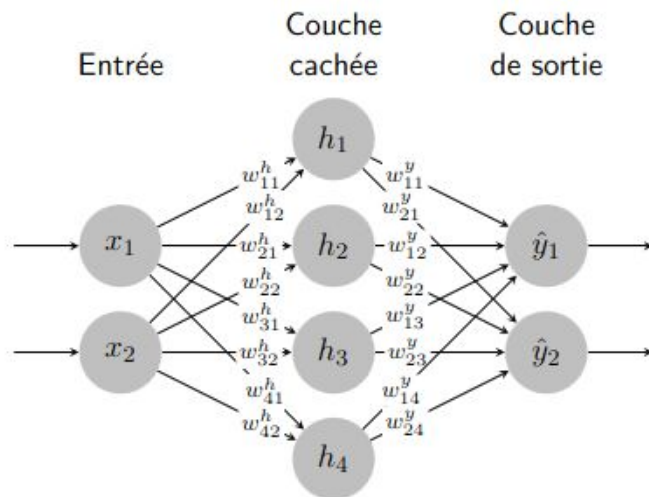
$$l_{CE}(\hat{y}, y) = - \sum_{i=0}^{\#Classes-1} y_i \log(\hat{y}_i)$$



$$\begin{cases} \tilde{h}_i = \sum_{j=1}^{n_x} W_{i,j}^h x_j + b_i^h \\ h_i = \tanh(\tilde{h}_i) \\ \tilde{y}_i = \sum_{j=1}^{n_h} W_{i,j}^y h_j + b_i^y \\ \hat{y}_i = \text{SoftMax}(\tilde{y}_i) = \frac{e^{\tilde{y}_i}}{\sum_{j=1}^{n_y} e^{\tilde{y}_j}} \end{cases}$$

Example: Tanh MLP

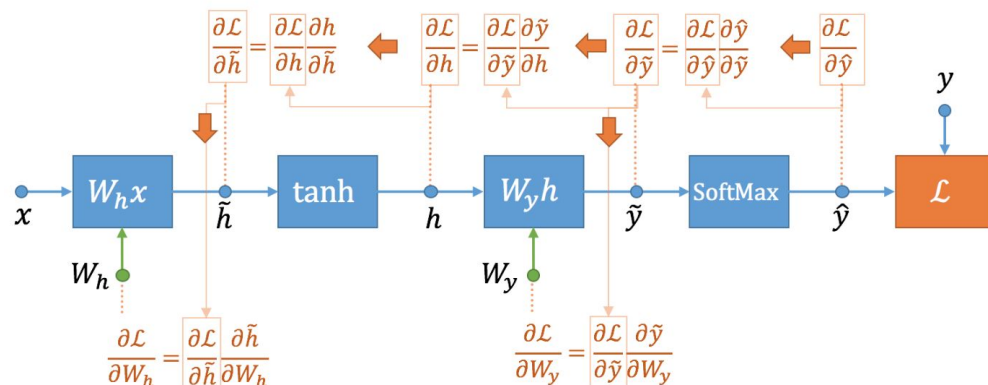
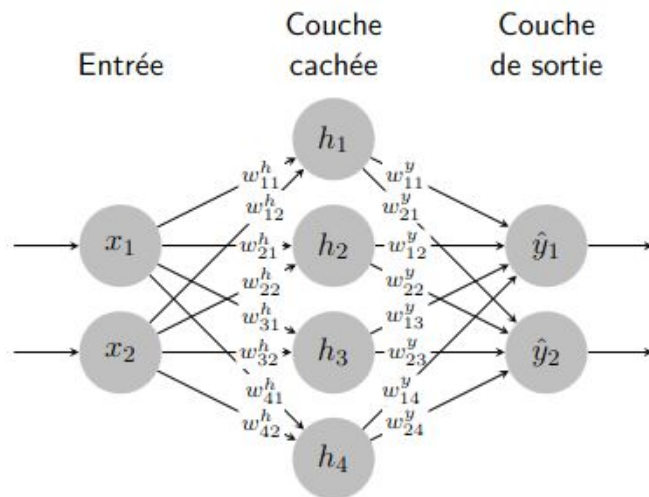
$$l_{CE}(\hat{y}, y) = - \sum_{i=0}^{\#Classes-1} y_i \log(\hat{y}_i)$$



$$\begin{cases} \tilde{h}_i = \sum_{j=1}^{n_x} W_{i,j}^h x_j + b_i^h \\ h_i = \tanh(\tilde{h}_i) \\ \tilde{y}_i = \sum_{j=1}^{n_h} W_{i,j}^y h_j + b_i^y \\ \hat{y}_i = \text{SoftMax}(\tilde{y}_i) = \frac{e^{\tilde{y}_i}}{\sum_{j=1}^{n_y} e^{\tilde{y}_j}} \end{cases}$$

Example: Tanh MLP


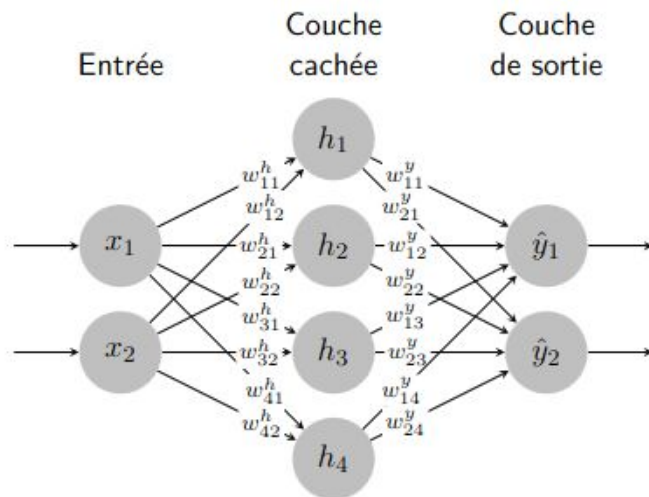
$$l_{CE}(\hat{y}, y) = - \sum_{i=0}^{\#Classes-1} y_i \log(\hat{y}_i)$$



$$\left\{ \begin{array}{l} \tilde{h}_i = \sum_{j=1}^{n_x} W_{i,j}^h x_j + b_i^h \\ h_i = \tanh(\tilde{h}_i) \\ \tilde{y}_i = \sum_{j=1}^{n_h} W_{i,j}^y h_j + b_i^y \\ \hat{y}_i = \text{SoftMax}(\tilde{y}_i) = \frac{e^{\tilde{y}_i}}{\sum_{j=1}^{n_y} e^{\tilde{y}_j}} \end{array} \right\} \quad \left\{ \begin{array}{l} \delta_i^y = \frac{\partial \ell}{\partial \tilde{y}_i} = \\ \frac{\partial \ell}{\partial W_{i,j}^y} = \\ \frac{\partial \ell}{\partial b_i^y} = \end{array} \right.$$

Example: Tanh MLP


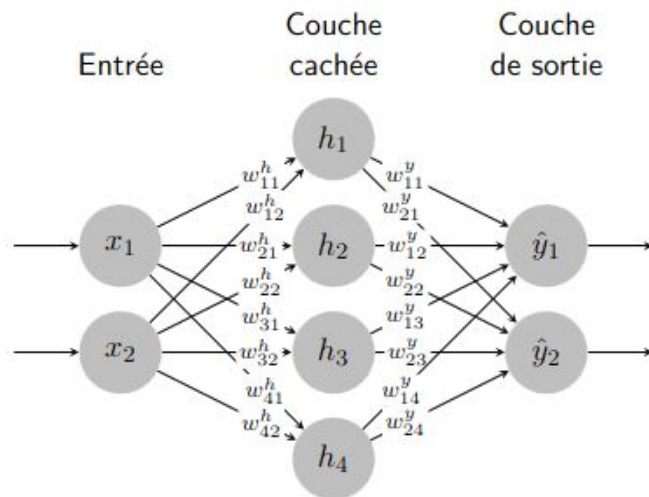
$$l_{CE}(\hat{y}, y) = - \sum_{i=0}^{\#Classes-1} y_i \log(\hat{y}_i)$$



$$\left\{ \begin{array}{l} \tilde{h}_i = \sum_{j=1}^{n_x} W_{i,j}^h x_j + b_i^h \\ h_i = \tanh(\tilde{h}_i) \\ \tilde{y}_i = \sum_{j=1}^{n_h} W_{i,j}^y h_j + b_i^y \\ \hat{y}_i = \text{SoftMax}(\tilde{y}_i) = \frac{e^{\tilde{y}_i}}{\sum_{j=1}^{n_y} e^{\tilde{y}_j}} \end{array} \right\} \quad \left\{ \begin{array}{l} \delta_i^y = \frac{\partial \ell}{\partial \tilde{y}_i} = \hat{y}_i - y_i \\ \frac{\partial \ell}{\partial W_{i,j}^y} = \delta_i^y h_j \\ \frac{\partial \ell}{\partial b_i^y} = \delta_i^y \end{array} \right.$$

Example: Tanh MLP


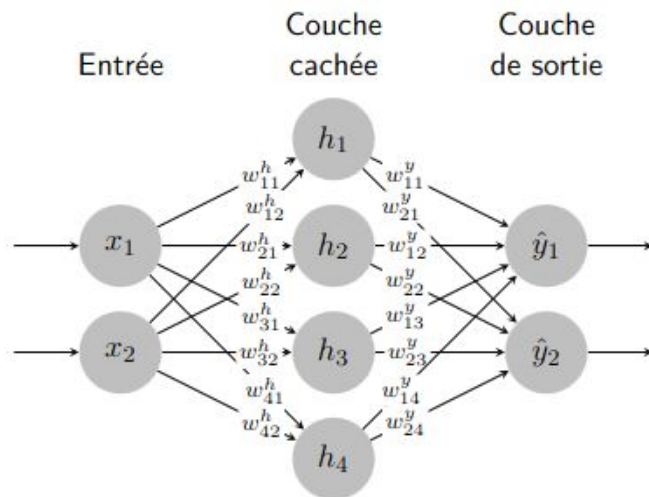
$$l_{CE}(\hat{y}, y) = - \sum_{i=0}^{\#Classes-1} y_i \log(\hat{y}_i)$$



$$\left\{ \begin{array}{l} \tilde{h}_i = \sum_{j=1}^{n_x} W_{i,j}^h x_j + b_i^h \\ h_i = \tanh(\tilde{h}_i) \\ \tilde{y}_i = \sum_{j=1}^{n_h} W_{i,j}^y h_j + b_i^y \\ \hat{y}_i = \text{SoftMax}(\tilde{y}_i) = \frac{e^{\tilde{y}_i}}{\sum_{j=1}^{n_y} e^{\tilde{y}_j}} \end{array} \right. \quad \left\{ \begin{array}{l} \delta_i^y = \frac{\partial \ell}{\partial \tilde{y}_i} = \hat{y}_i - y_i \\ \frac{\partial \ell}{\partial W_{i,j}^y} = \delta_i^y h_j \\ \frac{\partial \ell}{\partial b_i^y} = \delta_i^y \\ \delta_i^h = \frac{\partial \ell}{\partial \tilde{h}_i} = \\ \frac{\partial \ell}{\partial W_{i,j}^h} = \\ \frac{\partial \ell}{\partial b_i^h} = \end{array} \right.$$

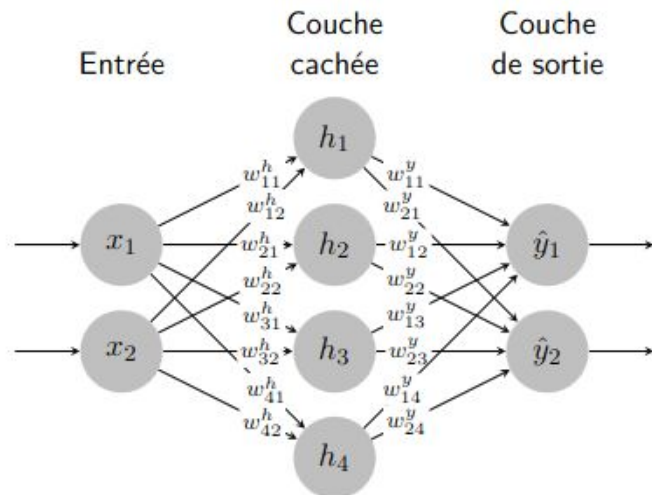
Example: Tanh MLP

$$l_{CE}(\hat{y}, y) = - \sum_{i=0}^{\#Classes-1} y_i \log(\hat{y}_i)$$



$$\left\{ \begin{array}{l} \tilde{h}_i = \sum_{j=1}^{n_x} W_{i,j}^h x_j + b_i^h \\ h_i = \tanh(\tilde{h}_i) \\ \tilde{y}_i = \sum_{j=1}^{n_h} W_{i,j}^y h_j + b_i^y \\ \hat{y}_i = \text{SoftMax}(\tilde{y}_i) = \frac{e^{\tilde{y}_i}}{\sum_{j=1}^{n_y} e^{\tilde{y}_j}} \end{array} \right. \quad \left\{ \begin{array}{l} \delta_i^y = \frac{\partial \ell}{\partial \tilde{y}_i} = \hat{y}_i - y_i \\ \frac{\partial \ell}{\partial W_{i,j}^y} = \delta_i^y h_j \\ \frac{\partial \ell}{\partial b_i^y} = \delta_i^y \\ \delta_i^h = \frac{\partial \ell}{\partial \tilde{h}_i} = (1 - h_i^2) \sum_{j=1}^{n_y} \delta_j^y W_{j,i}^y \\ \frac{\partial \ell}{\partial W_{i,j}^h} = \delta_i^h x_j \\ \frac{\partial \ell}{\partial b_i^h} = \delta_i^h \end{array} \right.$$

- Core problem: Find gradient updates
- Gradient can be computed efficiently with backpropagation
 - Chain rule starting from the “end”
 - Re-use computed gradients (Bottom-up DP)
 - Keep forward activations for gradients
 - Simple layers mean simple gradient blocks



$$\begin{cases} \tilde{\mathbf{h}} = \mathbf{x}\mathbf{W}^h + \mathbf{b}^h \\ \mathbf{h} = \tanh(\tilde{\mathbf{h}}) \\ \tilde{\mathbf{y}} = \mathbf{h}\mathbf{W}^y + \mathbf{b}^y \\ \hat{\mathbf{y}} = \text{SoftMax}(\tilde{\mathbf{y}}) \end{cases}$$

$$\begin{cases} \nabla_{\tilde{\mathbf{y}}} = \hat{\mathbf{y}} - \mathbf{y} \\ \nabla_{\mathbf{W}^y} = \nabla_{\tilde{\mathbf{y}}}^\top \mathbf{h} \\ \nabla_{\mathbf{b}^y} = \nabla_{\tilde{\mathbf{y}}}^\top \\ \nabla_{\tilde{\mathbf{h}}} = (\nabla_{\tilde{\mathbf{y}}} \mathbf{W}^y) \odot (1 - \mathbf{h}^2) \\ \nabla_{\mathbf{W}^h} = \nabla_{\tilde{\mathbf{h}}}^\top \mathbf{x} \\ \nabla_{\mathbf{b}^h} = \nabla_{\tilde{\mathbf{h}}}^\top \end{cases}$$

- Lab2 on Moodle
 - Implement this by hand with basic torch!
 - Careful with batch dimension!


```
def backward(params, outputs, Y):
    bsize = Y.shape[0]
    grads = {}

    Y_tilde_grad = outputs["yhat"] - Y
    h_tilde_grad = torch.mm(Y_tilde_grad, params["Wy"])
    h_tilde_grad = h_tilde_grad * (1 - torch.pow(outputs["h"], 2))

    grads["Wy"] = torch.mm(Y_tilde_grad.T, outputs["h"])
    grads["Wh"] = torch.mm(h_tilde_grad.T, outputs["X"])
    grads["by"] = Y_tilde_grad.sum(dim=0, keepdim=True).T
    grads["bh"] = h_tilde_grad.sum(0, keepdim=True).T

    grads["Wy"] /= bsize
    grads["by"] /= bsize
    grads["Wh"] /= bsize
    grads["bh"] /= bsize

    return grads
```

- Torch.tensor object
 - Np.array like
 - Tracked on a computational graph
 - .grad variable to track gradients
 - .backward to backpropagate gradients through the graph
 - Activate .autograd!

```
def backward(params, outputs, Y):
    bsize = Y.shape[0]
    grads = {}

    Y_tilde_grad = outputs["yhat"] - Y
    h_tilde_grad = torch.mm(Y_tilde_grad, params['Wy'])
                        * (1 - torch.pow(outputs['h'], 2))

    grads["Wy"] = torch.mm(Y_tilde_grad.T, outputs["h"])
    grads["Wh"] = torch.mm(h_tilde_grad.T, outputs["X"])
    grads["by"] = Y_tilde_grad.sum(dim=0, keepdim=True).T
    grads["bh"] = h_tilde_grad.sum(0, keepdim=True).T

    grads['Wy'] /= bsize
    grads['by'] /= bsize
    grads['Wh'] /= bsize
    grads['bh'] /= bsize

    return grads
```

```
params['Wh'] = torch.randn(nh, nx) * 0.3
params['Wh'].requires_grad = True
params['bh'] = torch.zeros(nh, 1, requires_grad=True)
params['Wy'] = torch.randn(ny, nh) * 0.3
params['Wy'].requires_grad = True
params['by'] = torch.zeros(ny, 1, requires_grad=True)
```

```
with torch.no_grad():
    params['Wy'] -= eta * params['Wy'].grad
    params['Wh'] -= eta * params['Wh'].grad
    params['by'] -= eta * params['by'].grad
    params['bh'] -= eta * params['bh'].grad

    params['Wy'].grad.zero_()
    params['Wh'].grad.zero_()
    params['by'].grad.zero_()
    params['bh'].grad.zero_()
```